

HYBRA: A Deductive Program Verifier for Hyperproperties

ANONYMOUS AUTHOR(S)

Hyperproperties relate multiple executions of a program and are useful to express common correctness properties (such as determinism) and security properties (such as non-interference). While there are a number of powerful program logics for the deductive verification of hyperproperties, their automation falls behind. Most existing deductive verification tools are limited to safety properties, but cannot reason about the existence of executions, for instance, to prove the violation of a safety property. Others support more flexible hyperproperties such as generalized non-interference, but have limitations in terms of the programs and proof structures they support.

In this paper, we present the first deductive verification technique for arbitrary hyperproperties over multiple executions of the same program. Our technique automates the generation of verification conditions for Hyper Hoare Logic. Our key insight is that arbitrary hyperproperties and the corresponding proof rules can be encoded into a standard intermediate verification language by representing *sets of states* of the input program *explicitly* in the states of the intermediate program. Verification is then automated using an existing SMT-based verifier for the intermediate language. We implement our technique in a tool called HYBRA and demonstrate that it can reliably verify complex hyperproperties.

CCS Concepts: • **Theory of computation** → **Logic and verification**; **Hoare logic**.

Additional Key Words and Phrases: Hyperproperties, Program Logic, Incorrectness Logic

1 INTRODUCTION

Hyperproperties [Clarkson and Schneider 2008] relate multiple executions of a program. They can be used to express correctness properties such as determinism (two executions of a program with the same input will result in the same output) or monotonicity (executing the program with a larger input will result in a larger output). Hyperproperties are also used to express security and information-flow properties, such as non-interference.

There are numerous program logics that enable the formal verification of different kinds of hyperproperties. Relational Hoare Logic (RHL) [Benton 2004] allows to reason about 2-safety hyperproperties, that is, properties that should hold *for all pairs of executions* of the same program (such as determinism). Since these properties relate all possible “first” executions to all possible “second” executions, we call them $\forall\forall$ -properties. RHL has then been extended to support k -safety hyperproperties [Sousa and Dillig 2016], i.e., properties that should hold *for all combinations of k executions* of the same program (\forall^k -properties), such as transitivity ($k = 3$) or associativity ($k = 4$). We call program logics for safety hyperproperties (i.e., \forall^* -properties) *overapproximate*, because they work by overapproximating the set of possible executions.

Many important hyperproperties fall outside the class of k -safety hyperproperties, because they require proving the *existence* of executions. For the special case of one execution, Reverse Hoare Logic [de Vries and Koutavas 2011] and Incorrectness Logic [O’Hearn 2019] allow one to prove the existence of an execution, for instance, the reachability of a bug. We call such properties \exists -properties; the corresponding logics *underapproximate* the set of possible executions.

More complex hyperproperties, of the form $\forall^*\exists^*$ or $\exists^*\forall^*$, require program logics that can perform both overapproximate and underapproximate reasoning. An important example is *Generalized Non-Interference* (GNI) [McCullough 1987; McLean 1996], a $\forall\forall\exists$ -hyperproperty. GNI requires, for any two executions τ_1 and τ_2 with the same low-sensitivity inputs, the existence of a third execution

with the same high inputs as τ_1 and the same low output as τ_2 . Examples of $\exists^*\forall^*$ -hyperproperties include the violation of GNI ($\exists\exists\forall$) and the existence of an execution with minimal values ($\exists\forall$). Several recent program logics support these more-complex hyperproperties. RHLE [Dickerson et al. 2022] supports $\forall^*\exists^*$ -hyperproperties, BiKAT [Antonopoulos et al. 2023] supports $\forall\exists$ -hyperproperties and (in principle) $\exists\forall$ -hyperproperties, and Hyper Hoare Logic [Dardinier and Müller 2023] supports hyperproperties with arbitrary quantifier alternation, including both $\forall^*\exists^*$ and $\exists^*\forall^*$ -hyperproperties.

Automation. Verification in program logics for safety properties of single executions (such as Hoare Logic) can be automated using *deductive program verifiers* (verifiers for short), such as BOOGIE [Leino 2008], DAFNY [Leino 2010], VIPER [Müller et al. 2016], or WHY3 [Filliâtre and Paskevich 2013]. Given a program, a specification, and hints from the user (such as loop invariants), these tools attempt to prove that the program satisfies the specification by computing proof obligations and solving them using an SMT solver.

Deductive verifiers for hyperproperties are mostly limited to k -safety hyperproperties, which can be reduced to safety properties for a product program [Barthe et al. 2011; Eilers et al. 2019] and then verified using an off-the-shelf verifier. Alternatively, there are dedicated verifiers for hyperproperties, such as WhyRel [Nagasamudram et al. 2023], SECC (based on SecCSL [Ernst and Murray 2019]), and HYPERVIPER (based on CommCSL [Eilers et al. 2023]) for 2-safety hyperproperties, and DESCARTES (based on Cartesian Hoare Logic [Sousa and Dillig 2016]) for k -safety hyperproperties.

ORHLE (based on RHLE [Dickerson et al. 2022]) is the only deductive verifier that goes beyond k -safety hyperproperties by supporting $\forall^*\exists^*$ -hyperproperties. However, it is limited to a *fixed* quantification scheme; users have to first fix the numbers of \forall -quantifiers and \exists -quantifiers and then write preconditions, postconditions, and loop invariants in this fixed scheme. It is, thus, not possible to compose proofs with different quantification schemes, e.g., to use a \forall -property in the proof of a $\forall\forall$ -property. Moreover, ORHLE supports relational loop invariants only when the different executions perform the same number of loop iterations, which limits the programs and hyperproperties that can be verified in practice.

To the best of our knowledge, no existing verifier supports properties beyond $\forall^*\exists^*$ -hyperproperties, such as $\exists^*\forall^*$ -hyperproperties.

This work. We present the first deductive verifier for hyperproperties with arbitrary quantifier alternations. Our tool, HYPRA, is based on Hyper Hoare Logic (HHL) [Dardinier and Müller 2023]. Like HHL, it allows assertions to quantify explicitly over states, giving users the flexibility to express and combine different types of hyperproperties in the same proof. Going beyond HHL, HYPRA supports reasoning about runtime errors (e.g., to prove the existence or absence of bugs).

Our key insight is that arbitrary hyperproperties and the corresponding proof rules can be encoded into a standard intermediate verification language by representing *sets of states* of the input program *explicitly* in the states of the intermediate program. Verification is then automated using an existing SMT-based verifier for the intermediate language. To ensure that SMT solvers can handle the resulting verification conditions, our encoding carefully manages quantifier instantiations by tracking simultaneously an over- and an underapproximation of the sets of states. Note that, in contrast to product constructions, our encoding does *not* duplicate the statements of the input program and can handle arbitrary hyperproperties (beyond k -safety).

Like HHL, we focus on hyperproperties that relate multiple executions of the *same* program; relating executions of *different* programs (e.g., to prove their equivalence) poses additional challenges (such as finding an alignment), which are orthogonal to the problems addressed here. Both WhyRel and ORHLE support such relational proofs.

Contributions. In summary, our work makes the following contributions:

- (1) We extend Hyper Hoare Logic [Dardinier and Müller 2023] with the ability to reason about runtime errors. This allows us to prove correctness (the absence of bugs), incorrectness (the existence of bugs) and more complex hyperproperties (e.g., proving that the occurrence of a runtime error does not leak any secret information).
- (2) We present the first approach to generate verification conditions for hyperproperties with arbitrary quantifier alternations for loop-free statements. The resulting verification conditions are amenable to SMT solvers.
- (3) HHL provides multiple loop rules for different kinds of programs and properties. We show how to select the right rule automatically, such that users are not exposed to the details of the underlying logic. This automatic selection is crucial when dealing with $\exists^*\forall^*$ loop invariants, which require the application of several different loop rules. Moreover, we present and prove sound a new loop rule for $\forall^*\exists^*$ -hyperproperties, which is easier to automate than the corresponding rule in HHL.
- (4) We implement our approach in a tool called HYPRA, based on the VIPER intermediate language [Müller et al. 2016]. Our evaluation on a set of benchmarks from the literature shows that HYPRA can prove a large class of hyperproperties for a large class of programs, in a reasonable amount of time and with a reasonable amount of proof annotation.

Outline. The rest of the paper is organized as follows: Sect. 2 highlights the capabilities of our verifier through several examples, and presents our extension of Hyper Hoare Logic to reason about runtime errors. Sect. 3 presents our approach to generate verification conditions for loop-free statements, while Sect. 4 handles loops. We discuss the implementation and evaluation of HYPRA in Sect. 5, related work in Sect. 6, and limitations and future work in Sect. 7.

2 A TOUR OF THE VERIFIER

In this section, we illustrate the key capabilities of our verifier on several examples. Sect. 2.1 shows how HYPRA can perform both over- and underapproximation, and how combining both allows us to prove complex hyperproperties (such as $\exists^*\forall^*$ -hyperproperties). Sect. 2.2 illustrates how HYPRA reasons about runtime errors, in particular how it can prove the absence of errors, the existence of errors, and more complex (hyper)properties (such as the fact that the occurrence of a runtime error does not leak secret information). We also explain how we extend Hyper Hoare Logic (which does not support errors) to do so. Finally, Sect. 2.3 shows how HYPRA handles while loops. All examples shown in this section are successfully and automatically verified by our tool.

2.1 Overapproximation, Underapproximation, and Hyperproperties

HYPRA establishes *hyper-triples* of the form $[P] C [Q]$, where C is a program statement, and P (the precondition) and Q (the postcondition) are *hyper-assertions*, i.e., predicates over sets of states (which we formally define in Sect. 2.2). Intuitively, the triple $[P] C [Q]$ means that for any set S of initial states that satisfies the precondition P , the set S' of all states (including error states) reachable by executing C from any state in S satisfies the postcondition Q .

Over- and underapproximation. As an example, consider the method `randNat` in Fig. 1 (left). This method non-deterministically chooses a value for x (which can for example represent a random choice or some user input), and assigns 1 to y if $x > 0$, and 2 otherwise. In other words, this method non-deterministically returns 1 or 2. The keyword **requires** specifies the precondition of the method, while the keyword **ensures** specifies the postcondition. Multiple preconditions or postconditions are simply conjoined. Thus, this example corresponds to the hyper-triple

$$[\exists\langle\sigma\rangle. \top] C_r [(\forall\langle\sigma\rangle. 1 \leq \sigma(y) \leq 2) \wedge (\exists\langle\sigma\rangle. \sigma(y) = 1) \wedge (\exists\langle\sigma\rangle. \sigma(y) = 2)]$$

```

148 method randNat() returns (y: Int)
149   requires  $\exists \langle \sigma \rangle. \top$ 
150   ensures  $\forall \langle \sigma \rangle. 1 \leq \sigma(y) \leq 2$ 
151   ensures  $\exists \langle \sigma \rangle. \sigma(y) = 1$ 
152   ensures  $\exists \langle \sigma \rangle. \sigma(y) = 2$ 
153 {
154   var x: Int
155   x := nonDet() // {hint}
156   // use hint(0,1)
157   if (x > 0) {
158     y := 1
159   }
160   else {
161     y := 2
162   }
163 }
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196

```

```

method secure(h: Int, l: Int) returns (o: Int)
  requires  $\forall \langle \sigma_1 \rangle, \langle \sigma_2 \rangle. \sigma_1(l) = \sigma_2(l)$ 
  ensures  $\forall \langle \sigma_1 \rangle, \langle \sigma_2 \rangle. \sigma_1(o) = \sigma_2(o)$ 
{
  if (h > 0) { o := 2 * l }
  else { o := l }
  if (h <= 0) { o := o + l }
}

```

```

method leaky(h: Int) returns (o: Int)
  requires  $\exists \langle \sigma_1 \rangle, \langle \sigma_2 \rangle. \sigma_1(h) < \sigma_2(h)$ 
  ensures  $\exists \langle \sigma_1 \rangle, \langle \sigma_2 \rangle. \forall \langle \sigma \rangle. \sigma(o) = \sigma_1(o) \Rightarrow \sigma(h) \neq \sigma_2(h)$ 
{
  var y: Int
  y := nonDet() // {hint}
  assume 0 <= y <= 10
  // use hint(10)
  o := h + y
}

```

Fig. 1. Examples of overapproximation, underapproximation, and hyperproperties. The example on the left illustrates \forall -properties and \exists -properties of individual program executions. The examples on the right illustrate hyperproperties. Method `secure` satisfies non-interference, a $\forall\forall$ -hyperproperty. Method `leaky` (which is adapted from Dardinier and Müller [2023]) violates generalized non-interference. The negation of this property is expressed as an $\exists\exists\forall$ -hyperproperty. All examples are successfully verified by HYPRA, using the provided hints to construct witnesses for existential quantifiers.

where C_r refers to the body of method `randNat`. Let S' be the set of reachable states at the end of the method. The postcondition $\forall \langle \sigma \rangle. 1 \leq \sigma(y) \leq 2$, which is equivalent to $\forall \sigma \in S'. 1 \leq \sigma(y) \leq 2$, *overapproximates* the set S' ; It means that, in any final state (from S'), y will either be 1 or 2, corresponding to the standard Hoare triple $\{\top\} C_r \{1 \leq y \leq 2\}$. On the other hand, the postconditions $\exists \langle \sigma \rangle. \sigma(y) = 1$ and $\exists \langle \sigma \rangle. \sigma(y) = 2$, equivalent to $\exists \sigma \in S'. \sigma(y) = 1$ and $\exists \sigma \in S'. \sigma(y) = 2$, respectively, *underapproximate* the set S' : They express the existence of two reachable final states with $y = 1$ and $y = 2$. Conjoined, these three postconditions express that this method has only two possible outcomes for y , 1 and 2, and that both outcomes are reachable. The precondition $\exists \langle \sigma \rangle. \top$ expresses that the set of initial states is non-empty. This precondition is required for the hyper-triple to hold, otherwise the postconditions $\exists \langle \sigma \rangle. \sigma(y) = 1$ and $\exists \langle \sigma \rangle. \sigma(y) = 2$ would not hold (because no states are reachable from an empty set of initial states).

While the first postcondition verifies automatically, proving the existentially-quantified postconditions requires a user-provided hint. *Hints* are annotations for non-deterministic assignments that give examples of values that might be assigned. HYPRA uses this information to construct witness states for \exists -properties. In our example, the `{hint}` annotation after the non-deterministic assignment introduces an identifier for this assignment, and the annotation `use hint(0,1)` tells the verifier that 0 and 1 are relevant choices for this non-deterministic assignment (technically, hints are used to provide triggers for the quantifier instantiation in the SMT solver).¹ The two

¹Note that, in general, hints can depend on the variables of one or more states. For example, `use $\forall \langle \sigma \rangle. \text{hint}(\sigma(n))$` tells the SMT solver to consider the value of variable n in all relevant states σ . Hints can also depend on multiple quantified states, as in `use $\forall \langle \sigma_1 \rangle, \langle \sigma_2 \rangle. \text{hint}(\sigma_1(a) + \sigma_2(b))$` .

values ensure that both branches of the subsequent conditional statement are considered, which is necessary to prove both existentially-quantified postconditions. The specific values are irrelevant in this example; any pair of a non-negative and a positive integer would work.

Hyperproperties. Overapproximating the set of reachable states allows us to formally verify safety hyperproperties such as non-interference, as illustrated by method `secure` in Fig. 1 (top right). The specification expresses that the output o depends only on the low-sensitive input l and, thus, does not leak any information about the secret input h . HYPPRA verifies this example without further annotations.

By both overapproximating and underapproximating the set of reachable states, our approach can verify more complex hyperproperties, such as $\forall^*\exists^*$ -hyperproperties or $\exists^*\forall^*$ -hyperproperties, as illustrated by the method `leaky` in Fig. 1 (bottom right). The statements $y := \text{nonDet}()$ and `assume $0 \leq y \leq 10$` model a non-deterministic choice between 0 and 10. This method leaks information about its secret input h via its output o : h is between $o - 10$ and o . To prove this claim, we formally verify that the method violates generalized non-interference [McCullough 1987; Volpano et al. 1996]. That is, we prove the existence of two executions with distinct secret values for h that can be *distinguished*. The postcondition expresses that observing the output $\sigma_1(o)$ rules out the possibility that the secret value of h was $\sigma_2(h)$, thus leaking information about the initial value of h . Verifying this existentially-quantified postcondition requires a hint; choosing the provided value 10 for $\sigma_2(y)$ yields the required witnesses.

2.2 Reasoning about Runtime Errors

Our examples so far reason about properties of *normal states*, that is, states that are reached when the program executes successfully. In addition, our technique can also reason about a set of *error states*, which are reached when a runtime error occurs. This features allows us to prove both the absence and presence of runtime errors, as well as advanced hyperproperties such as failure-sensitive non-interference.

Method `buggy` in Fig. 2 (top left) calls method `randNat` (see Fig. 1), doubles the result, and asserts that the resulting value is odd. The postconditions express the existence of two failing executions, à la Incorrectness Logic [O'Hearn 2019]: There exist error states $\langle \sigma \rangle_{\text{er}}$ where the result of `randNat` is 1 and 2, respectively.

The ability to quantify over error states allows us to express more complex properties. For example, the specification of method `almostCorrect` in Fig. 2 (top right) expresses that a runtime error occurs *only if* the non-deterministic value assigned to o is 0 *and* the input x is 0. This *almost-correctness* is captured by the universal quantification in the postcondition which expresses that all error states satisfy $x = 0 \wedge o = 0$, that is, no other execution fails.

The absence of errors can be specified via the postcondition $\forall \langle \sigma \rangle_{\text{er}}. \perp$, which expresses that the set of error states is empty.

In the context of hyperproperties, reasoning about error states is, for instance, useful to express failure-sensitive non-interference, that is, the fact that observing a runtime error does not leak secret information. The specification of method `lowError` in Fig. 2 (bottom) expresses this property. For two different executions with the same value for the low-sensitive input l (but potentially different values for the secret input h), one execution fails if and only if the other execution fails. In particular, this proves that the occurrence of a runtime error does not depend on h , such that observing an error does not leak secret information. In this specification, the parameter t is used to *tag* executions, which allows us to identify the pre- and post-state of a given execution. Tags are expressed as logical variables in Hyper Hoare Logic, but represented by (immutable) program variables in HYPPRA.

```

246 method buggy() returns (x: Int)
247   requires  $\exists \langle \sigma \rangle. \top$ 
248   ensures  $\exists \langle \sigma \rangle_{er}. \sigma(x) = 1$ 
249   ensures  $\exists \langle \sigma \rangle_{er}. \sigma(x) = 2$ 
250 {
251   x := randNat()
252   var y: Int := x + x
253   assert y % 2 == 1
254 }

```

```

method almostCorrect(x: Int) returns (o: Int)
  requires  $\forall \langle \sigma \rangle. \sigma(x) \geq 0$ 
  ensures  $\forall \langle \sigma \rangle_{er}. \sigma(x) = 0 \wedge \sigma(o) = 0$ 
  {
    o := nonDet()
    assume o >= 0
    var y: Int := x + o
    assert y > 0
  }

```

```

method lowError(h: Int, l: Int, t: Int)
  requires  $\exists \langle \sigma_1 \rangle, \langle \sigma_2 \rangle. \sigma_1(t) = 1 \wedge \sigma_2(t) = 2$ 
  requires  $\forall \langle \sigma_1 \rangle, \langle \sigma_2 \rangle. \sigma_1(l) = \sigma_2(l)$ 
  ensures  $(\exists \langle \sigma_1 \rangle_{er}. \sigma_1(t) = 1) \Leftrightarrow (\exists \langle \sigma_2 \rangle_{er}. \sigma_2(t) = 2)$ 
  {
    if (h > 0) {
      assert l >= 0
    }
    if (l < 0) {
      assert false
    }
  }

```

Fig. 2. Reasoning about runtime errors. In the top left example, the postconditions specify two possible executions that lead to a runtime error (an assertion violation). The postcondition on the top right expresses that the method fails *only if* both x and o are 0. The example on the bottom illustrates reasoning about runtime errors in the context of hyperproperties. The specification expresses that the occurrence of a runtime error does not depend on the secret input h. All examples are successfully verified by HYPRA without any hints.

Extending Hyper Hoare Logic to support runtime errors. We have extended Hyper Hoare Logic, which does not provide any support for reasoning about errors, as follows:

DEFINITION 1. Hyper-triples with errors.

Given a program statement C and two program states σ and σ' , we write $\langle C, \sigma \rangle \rightarrow \sigma'$ to express that executing C in the initial state σ may lead to the final state σ' . Executions that lead to runtime errors (because of violated assertions) are denoted as $\langle C, \sigma \rangle_{er} \rightarrow \sigma'$, where σ' is the last state reached before the error occurred. We refer to such states as error states, in contrast to normal states. The set of normal states reachable by executing C in any state from S is denoted as $\text{sem}(C, S) \triangleq \{\sigma' \mid \exists \sigma \in S. \langle C, \sigma \rangle \rightarrow \sigma'\}$, while the set of error states reachable by executing C in any state from S is denoted as $\text{err}(C, S) \triangleq \{\sigma' \mid \exists \sigma \in S. \langle C, \sigma \rangle_{er} \rightarrow \sigma'\}$.

Hyper-assertions are predicates over pairs of sets of states, where the first set corresponds to normal states, and the second set corresponds to error states. Given two hyper-assertions P and Q , we write $\models [P] C [Q]$ to express that the triple $[P] C [Q]$ is valid, which is defined as follows:

$$\models [P] C [Q] \quad \text{iff} \quad \forall S. P(S, \emptyset) \Rightarrow Q(\text{sem}(C, S), \text{err}(C, S))$$

Note that we start with an *empty* set of error states (second argument of P). The reason is that, in the context of a sequential composition $C_1; C_2$, the set of error states that come from C_2 depends on $\text{sem}(C_1, S)$ only, but not on $\text{err}(C_1, S)$; formally, $\text{err}(C_1; C_2, S) = \text{err}(C_1, S) \cup \text{err}(C_2, \text{sem}(C_1, S))$. Thus, there would be no advantage in starting with a non-empty set of error states (i.e., no increased

```

295 method minimum(n: Int) returns (x: Int, y: Int)
296   requires  $\exists \langle \sigma \rangle. \top$ 
297   requires  $\forall \langle \sigma_1 \rangle, \langle \sigma_2 \rangle. \sigma_1(n) = \sigma_2(n)$ 
298   ensures  $\exists \langle \sigma \rangle. \forall \langle \sigma' \rangle. \sigma(x) \leq \sigma'(x) \wedge \sigma(y) \leq \sigma'(y)$ 
299 {
300   var i, r: Int
301   i, x, y := 0
302   while (i < n)
303     invariant  $\forall \langle \sigma_1 \rangle, \langle \sigma_2 \rangle. \sigma_1(n) = \sigma_2(n) \wedge \sigma_1(i) = \sigma_2(i)$ 
304     invariant  $\exists \langle \sigma \rangle. \forall \langle \sigma' \rangle. \sigma(x) \leq \sigma'(x) \wedge \sigma(y) \leq \sigma'(y)$ 
305     decreases n - i
306   {
307     r := nonDet() // {hint}
308     assume r >= 5
309     // use hint(5)
310     x := x + y + 2 * i + 3 * r
311     y := x + 3 * i + 2 * r
312     if (x >= n) { y := y + r }
313     i := i + 1
314   }
315 }
```

Fig. 3. Reasoning about loops. Given a loop invariant and an optional variant, HYPR automatically selects the appropriate loop rule. Like all other assertions, loop invariants may express arbitrary hyperproperties, here, the existence of an execution with minimal values for *x* and *y*. The example is successfully verified by HYPR.

expressivity), but there would be disadvantages: For example, to express that a method is partially correct, we would need to require the initial set of error states to be empty, which would prevent this method from being called in a context where errors have already happened.

Specification language. HYPR supports the following syntax for hyper-assertions P (which we used in the examples above), where E ranges over integer expressions, B over Boolean expressions, and P over hyper-assertions:

$$E ::= \sigma(y) \mid x \mid n \mid E + E \mid E - E \mid E * E \mid E / E \mid E \% E \mid \dots$$

$$B ::= \top \mid \perp \mid E = E \mid E > E \mid E \geq E \mid \neg B \mid \dots$$

$$P ::= \forall \langle \sigma \rangle. P \mid \exists \langle \sigma \rangle. P \mid \forall \langle \sigma \rangle_{\text{er}}. P \mid \exists \langle \sigma \rangle_{\text{er}}. P \mid \forall x. P \mid \exists x. P \mid P \wedge P \mid P \vee P \mid P \Rightarrow P \mid B$$

Note that the interaction with the set of normal states is restricted to the quantifiers $\forall \langle \sigma \rangle$ and $\exists \langle \sigma \rangle$, while the interaction with the set of error states is restricted to the quantifiers $\forall \langle \sigma \rangle_{\text{er}}$ and $\exists \langle \sigma \rangle_{\text{er}}$. Moreover, the quantifiers $\forall \langle \sigma \rangle_{\text{er}}$ and $\exists \langle \sigma \rangle_{\text{er}}$ are not allowed in preconditions (since we always start with an empty set of error states).

2.3 Reasoning about Loops

Hyper Hoare Logic provides four different loop rules that can prove different flavors of hyperproperties. These rules are applicable in different contexts; for example, some rules are applicable only if all loop executions perform the same number of iterations, and others only if the loop is proved to terminate. Based on a user-provided loop invariant and an optional loop variant, HYPR determines automatically which rule to apply. This allows users to reason about loops in a familiar way without

being exposed to the complexity of the underlying logic, as we will illustrate on method `minimum` in Fig. 3.

This method starts with $x = y = 0$ and performs n loop iterations, during which it modifies the values of x and y in a non-deterministic way. We want to prove that, given a fixed value for the input n (enforced by the precondition $\forall \langle \sigma_1 \rangle, \langle \sigma_2 \rangle. \sigma_1(n) = \sigma_2(n)$), there exists an execution where both x and y have minimal values at the end of the method (without specifying their values, which depend on n non-deterministic choices). The proof is based on a user-provided relational *loop invariant*, which must hold before the loop, and after every iteration. The first part of the loop invariant, $\forall \langle \sigma_1 \rangle, \langle \sigma_2 \rangle. \sigma_1(n) = \sigma_2(n) \wedge \sigma_1(i) = \sigma_2(i)$, ensures that all states have the same value for i and n , and thus that all executions will exit the loop simultaneously. Our verifier automatically detects this pattern, and uses a specialized loop encoding to handle it, as we will explain in Sect. 4. Moreover, knowing that all executions have the same value for i is necessary to prove the existence of an execution with minimal values. The second part of the loop invariant, $\exists \langle \sigma \rangle. \forall \langle \sigma' \rangle. \sigma(x) \leq \sigma'(x) \wedge \sigma(y) \leq \sigma'(y)$, ensures that, after any number of iterations, there exists an execution with minimal values for x and y , which corresponds to our postcondition. Finally, we need to prove that the loop terminates, otherwise our postcondition would not hold: If the loop did not terminate, then *no* final state would exist. We prove termination using a standard loop variant (following the **decreases** keyword). The verifier checks, for all states, that this loop variant is well-founded (non-negative), and that it strictly decreases during each iteration, thus ensuring termination.

This section illustrated the capabilities of our verification approach from a user's perspective. In the next two sections, we will explain how we compute verification conditions by encoding the input program into an intermediate verification language, for which an automated verifier exists.

3 VERIFICATION CONDITIONS FOR LOOP-FREE STATEMENTS

Given a loop-free program statement C (potentially containing hints), a precondition P and a postcondition Q , our verifier generates a Viper [Müller et al. 2016] program, such that validity of the Viper program implies validity of the hyper-triple $\models [P] C [Q]$. Our key insight is to represent *sets of states* of the input program C as *single states* in the Viper program. More precisely, the Viper program contains set-valued variables, whose contents represent the set of states. The generated Viper program starts with a set-valued variable S (containing an arbitrary value), assumes that S satisfies the precondition P (via an `assume`-statement in the Viper program), tracks the sets of normal states and error states that can be reached by executing C in any state from S (by updating the set-valued variable S accordingly), and checks whether they satisfy the postcondition Q (via an `assert`-statement in the Viper program). To simplify the presentation, we ignore the set of error states in the rest of this paper, and focus only on the set of normal states (which we also call the set of reachable states), but the same principles apply to both.

More precisely, our encoding tracks an *overapproximation* and an *underapproximation* of the set of reachable states. The overapproximation is sufficient to verify safety hyperproperties (\forall^* -hyperproperties), while the underapproximation is sufficient to verify \exists^* -hyperproperties. Reasoning about hyperproperties with arbitrary quantifier alternations, such as $\forall^* \exists^*$ or $\exists^* \forall^*$ -hyperproperties, requires combining both approximations. However, naively combining the two approximations does not work, as it can lead to matching loops where the SMT solver gets stuck in an infinite instantiation of quantifiers.

In Sect. 3.1, we present our approach to track an overapproximation and an underapproximation of the set of reachable states. We then explain, in Sect. 3.2, why naively combining these two

393	<code>// y := nonDet()</code>	<code>// y := nonDet()</code>
394	<code>assume $\forall \sigma_1 \in S_V^1. \exists \sigma_0, v. \sigma_0 \in S_V^0 \wedge \sigma_1 = \sigma_0 [y := v]$</code>	<code>assume $\forall \sigma_0, v. \sigma_0 \in S_\exists^0 \Rightarrow \sigma_0 [y := v] \in S_\exists^1$</code>
395	<code>// assume $0 \leq y \leq 10$</code>	<code>// assume $0 \leq y \leq 10$</code>
396	<code>assume $\forall \sigma_2 \in S_V^2. \sigma_2 \in S_V^1 \wedge 0 \leq \sigma_2(y) \leq 10$</code>	<code>assume $\forall \sigma_1 \in S_\exists^1. 0 \leq \sigma_1(y) \leq 10 \Rightarrow \sigma_1 \in S_\exists^2$</code>
397	<code>// o := h + y</code>	<code>// o := h + y</code>
398	<code>assume $\forall \sigma_3 \in S_V^3. \exists \sigma_2 \in S_V^2. \sigma_3 = \sigma_2 [o := \sigma_2(h) + \sigma_2(y)]$</code>	<code>assume $\forall \sigma_2 \in S_\exists^2. \sigma_2 [o := \sigma_2(h) + \sigma_2(y)] \in S_\exists^3$</code>
399		

Fig. 4. Viper encodings to compute the overapproximation (on the left) and underapproximation (on the right) of the set of reachable states, for the body of method `leaky` from Fig. 1. $S_V^0, S_V^1, S_V^2, S_V^3, S_\exists^0, S_\exists^1, S_\exists^2$, and S_\exists^3 are fresh variables. Additionally, S_V^0 and S_\exists^0 are assumed to satisfy together the precondition of the method (as we explain in Sect. 3.2), and S_V^3 and S_\exists^3 are the sets used to check whether the postcondition holds. $\sigma[x := v]$ denotes the state σ updated with x set to v .

approximations can lead to matching loops, and present our approach to combine these two approximations, which allows us to reason about hyperproperties with arbitrary quantifier alternations, while avoiding matching loops.

3.1 Overapproximating and Underapproximating the Set of Reachable States

Recall that the set of states reachable by executing a command C in any state from a set of initial states S is defined as $\text{sem}(C, S) \triangleq \{\sigma' \mid \exists \sigma \in S. \langle C, \sigma \rangle \rightarrow \sigma'\}$ (Def. 1). Equivalently, this definition can be seen as the conjunction of the two following properties: (1) every state $\sigma' \in \text{sem}(C, S)$ results from executing C in some state $\sigma \in S$, and (2) any final state σ' that results from executing C in a state $\sigma \in S$ belongs to $\text{sem}(C, S)$. More formally:

$$\forall \sigma'. \sigma' \in \text{sem}(C, S) \Rightarrow \exists \sigma. \sigma \in S \wedge (\langle C, \sigma \rangle \rightarrow \sigma') \quad (1)$$

$$\forall \sigma, \sigma'. \sigma \in S \wedge (\langle C, \sigma \rangle \rightarrow \sigma') \Rightarrow \sigma' \in \text{sem}(C, S) \quad (2)$$

Property (1) corresponds to an *upper bound* (i.e., an overapproximation) of the set of reachable states, while property (2) corresponds to a *lower bound* (i.e., an underapproximation). Leveraging those properties, our encoding tracks two set-valued variables S_V and S_\exists that respectively overapproximate and underapproximate the set of reachable states. We generate Viper encodings of the following form, where $\langle C, \sigma \rangle \rightarrow \sigma'$ is specialized for each atomic statement C (e.g., assignment, assert-statements, assume-statements) of the language, to update the current overapproximation S_V^n to the next overapproximation S_V^{n+1} , and the current underapproximation S_\exists^n to the next underapproximation S_\exists^{n+1} , for a program statement C :²

```

assume  $\forall \sigma_{n+1}. \sigma_{n+1} \in S_V^{n+1} \Rightarrow \exists \sigma_n. \sigma_n \in S_V^n \wedge \langle C, \sigma_n \rangle \rightarrow \sigma_{n+1}$  // update overapproximation
assume  $\forall \sigma_n, \sigma_{n+1}. \sigma_n \in S_\exists^n \wedge \langle C, \sigma_n \rangle \rightarrow \sigma_{n+1} \Rightarrow \sigma_{n+1} \in S_\exists^{n+1}$  // update underapproximation

```

Starting with $S_\exists^0 \subseteq S \subseteq S_V^0$, these encodings compute new values for S_V^{n+1} and S_\exists^{n+1} such that $S_\exists^{n+1} \subseteq \text{sem}(C, S) \subseteq S_V^{n+1}$. Note that the overapproximation S_V is sufficient to verify \forall^* -hyperproperties (safety hyperproperties), while the underapproximation S_\exists is sufficient to verify \exists^* -hyperproperties. Our tool uses this observation to optimize the encoding when only one kind of reasoning is needed.

²In practice, we have four variables: S_V and S_\exists , which represent the current approximations, and S'_V and S'_\exists , which represent the next approximations. Our encoding is actually of the form **havoc** S ; **assume** ...; $S := S'$. We use the superscripts n and $n+1$ to denote the values of these variables at different points in the encoding, to simplify the explanations.

Fig. 4 shows the concrete encodings generated with our approach for the method `leaky` from Fig. 1. The overapproximation encoding, on the left, can be read bottom-up. For any state σ_3 in S_V^3 at the end of the method, we learn that there exists a state $\sigma_2 \in S_V^2$ such that $\sigma_3 = \sigma_2[h := \sigma_2(h) + \sigma_2(y)]$. We then learn that this state σ_2 also belongs to S_V^1 , and satisfies $0 \leq \sigma_2(y) \leq 10$. Finally, we learn that there exist a state $\sigma_0 \in S_V^0$ and a value v such that $\sigma_2 = \sigma_0[y := v]$.

In contrast, the underapproximation encoding, on the right, should be read top-down. Starting with any state σ_0 in S_\exists^0 and any value (typically provided via a hint), we obtain that $\sigma_0[y := v]$ belongs to S_\exists^1 . If we can prove that $0 \leq v \leq 10$, we then obtain that σ_1 belongs to S_\exists^2 . Finally, we obtain that $\sigma_2[o := \sigma_2(h) + \sigma_2(y)]$ belongs to S_\exists^3 , which can be used as a witness to prove the postcondition of the method.

For conditional statements (not shown in Fig. 4), we leverage the fact that

$$\text{sem}(\text{if } (b) \{C_1\} \text{ else } \{C_2\}, S) = \text{sem}(\text{assume } b; C_1, S) \cup \text{sem}(\text{assume } \neg b; C_2, S)$$

Thus, to construct the overapproximation and underapproximation encodings for a conditional statement, we first construct S_V and S_\exists for the two branches separately, and then construct their unions.

3.2 Combining Overapproximation and Underapproximation

To reason about $\forall^* \exists^*$ and $\exists^* \forall^*$ -hyperproperties, we need to combine the two types of encodings, to both overapproximate (for universally-quantified states) and underapproximate (for existentially-quantified states) the set of reachable states. However, as we show below, naively combining the two approximations can lead to matching loops, where the SMT solver gets stuck in an infinite instantiation of quantifiers, in two different ways.

3.2.1 Naively combining the overapproximation and underapproximation encodings. The most straightforward way to combine the two approximations would be to assume that $S_V^n = S_\exists^n$ at every point of the encoding (i.e., for every n). Unfortunately, this quickly leads to matching loops, because of the interaction between the overapproximation and underapproximation encodings. As an example, consider the following overapproximation and underapproximation encodings for a non-deterministic assignment $y := \text{nonDet}()$ where $S^0 = S_V^0 = S_\exists^0$ and $S^1 = S_V^1 = S_\exists^1$.

```

assume  $\forall \sigma_1 \in S^1. \exists \sigma_0, v. \sigma_0 \in S^0 \wedge \sigma_1 = \sigma_0[y := v]$  // (overapproximation)
assume  $\forall \sigma_0, v. \sigma_0 \in S^0 \Rightarrow \sigma_0[y := v] \in S^1$  // (underapproximation)

```

At the level of the SMT solver, for any state $\sigma_1 \in S^1$, the overapproximation encoding introduces a new state $\sigma_1[y := v_0] \in S^0$ (for some value v_0). This subsequently triggers the universal quantifiers in the second **assume** statement, which proves (using $v = v_0$) that $\sigma_1[y := v_0] \in S^1$. This, in turn, triggers the universal quantifier in the first **assume** statement, which introduces a new state $\sigma_1[y := v_0][y := v_1] \in S^0$ (for some value v_1), which is then proven to be in S^1 by the underapproximation encoding, and so on, which results in an infinite cycle of quantifier instantiation.

To prevent this matching loop, we keep S_V^n and S_\exists^n as two different sets of states except when they both represent the set of *initial* states (as we motivate below), i.e., we assume (via an **assume**-statement in the Viper program) that $S^0 = S_V^0 = S_\exists^0$ at the beginning of any method, but we do not assume any relation between S_V^n and S_\exists^n for $n > 0$. In this way, any new state introduced via the \exists -quantifier in one encoding cannot trigger the instantiation of the \forall -quantifier in the other encoding, which circumvents matching loops.

To see why we want $S_V^0 = S_\exists^0$ at the beginning of the method, consider the method `simple` in Fig. 5, with its encoding shown on the right. The precondition tells us that for any state σ_1 , there exists a state σ_2 such that $\sigma_2(x) > \sigma_1(x)$, and that all states agree on the value of y . Thus, for any state σ_1 ,

```

491 method simple(x,y:Int) returns (z:Int) ...
492   requires  $\forall \langle \sigma_1 \rangle. \exists \langle \sigma_2 \rangle. \sigma_2(x) > \sigma_1(x)$  // next line is required to verify the program
493   requires  $\forall \langle \sigma_1 \rangle, \langle \sigma_2 \rangle. \sigma_1(y) = \sigma_2(y)$  assume  $S_V^0 = S_\exists^0$ 
494   ensures  $\forall \langle \sigma_1 \rangle. \exists \langle \sigma_2 \rangle. \sigma_2(z) > \sigma_1(z)$  assume  $\forall \sigma_1. \sigma_1 \in S_V^0 \Rightarrow \exists \sigma_2. \sigma_2 \in S_\exists^0 \wedge \sigma_2(x) > \sigma_1(x)$ 
495   { assume  $\forall \sigma_1, \sigma_2. \sigma_1 \in S_V^0 \wedge \sigma_2 \in S_\exists^0 \Rightarrow \sigma_1(y) = \sigma_2(y)$ 
496     z := x + y ... // encoding of z := x + y
497   } assert  $\forall \sigma_1. \sigma_1 \in S_V^1 \Rightarrow \exists \sigma_2. \sigma_2 \in S_\exists^1 \wedge \sigma_2(z) > \sigma_1(z)$ 

```

Fig. 5. A simple example that requires both overapproximation and underapproximation reasoning on the left, and its encoding on the right.




Fig. 6. Representation of which quantifiers can instantiate which other quantifiers. Each node represents the shape of a possible part of the precondition; For simplicity, we ignore shapes with more than two quantifiers. The grey arrow shows how instantiating the \forall^+ -quantifier in a $\forall^+ \exists^+$ -hyperproperty leads to the introduction of new existentially-quantified states, which can subsequently be used to instantiate universal quantifiers. The blue arrows show the existentially-quantified states that can be used to instantiate the universal quantifiers, and the red arrow shows the instantiations we forbid.

there should exist a state σ_2 such that $\sigma_2(z) > \sigma_1(z)$ after the assignment. However, without the assumption that $S_V^0 = S_\exists^0$ at the beginning of the method, this program would not be verified, as σ_2 belongs to S_\exists^0 , but not necessarily to S_V^0 , and thus one cannot prove that $\sigma_1(y) = \sigma_2(y)$.

3.2.2 Restricting quantifier instantiations for $\forall^* \exists^*$ -preconditions. Assuming that $S_V^0 = S_\exists^0$ at the beginning of the program is, however, not sufficient to rule out all matching loops, when the precondition contains a $\forall^* \exists^*$ -hyperproperty. One example is the first precondition of the method in Fig. 5, which is interpreted as $\forall \sigma_1 \in S_V^0. \exists \sigma_2 \in S_\exists^0. \sigma_2(x) > \sigma_1(x)$. The \forall -quantifier in the assertion introduces a new state $\sigma_2 \in S_\exists^0$ via the nested \exists -quantifier. Since we assume $S_V^0 = S_\exists^0$, the new state σ_2 can trigger the instantiation of the same \forall -quantifier, which in turn can introduce a new state $\sigma'_2 \in S_\exists^0$, and so on, leading to a matching loop.

Our solution is to use *limited* and *unlimited* functions [Leino and Monahan 2009] to control quantifier instantiations, to allow as many existentially-quantified states as possible to instantiate universal quantifiers, while avoiding matching loops. Concretely, when our tool translates a hyperassertion with a universal state-quantifier such as $\forall \langle \sigma \rangle. P$ into Viper, it checks whether P contains an existential state-quantifier: If so, then the \forall -quantifier is encoded with the *most restrictive* trigger (syntactic pattern), so that it can only be instantiated by existentially-quantified states that do not occur under a universal quantifier. Otherwise, the \forall -quantifier is encoded with a more permissive trigger, allowing it to be instantiated by all existentially-quantified states.

The effect of our solution is represented visually on Fig. 6. Each node (\exists^+ , $\exists^+ \forall^+$, $\forall^+ \exists^+$, \forall^+) represents the shape of a possible part of the precondition (we ignore shapes with more than two quantifiers for simplicity). The blue arrows show the instantiations *allowed* by our tool, while the

$$\begin{array}{c}
\frac{I \models \text{low}(b) \quad \models [I \wedge \Box b] C [I]}{\models [I] \text{ while } (b) \{C\} [(I \vee \Box \perp) \wedge \Box (\neg b)]} \text{ (WhileSync)} \\
\\
\frac{I \models \text{low}(b) \quad \models_{\parallel} [I \wedge \Box (b \wedge e = t)] C [I \wedge \Box (e < t)] \quad < \text{well-founded} \quad t \notin \text{fv}(I) \cup \text{mod}(C)}{\models_{\parallel} [I] \text{ while } (b) \{C\} [I \wedge \Box (\neg b)]} \text{ (WhileSyncTerm)} \\
\\
\frac{\models [I] \text{ if } (b) \{C\} [I] \quad \models [I] \text{ assume } \neg b [Q] \quad \text{no } \forall (_) \text{ after any } \exists \text{ in } Q}{\models [I] \text{ while } (b) \{C\} [Q]} \text{ (While-}\forall^* \exists^*) \\
\\
\frac{\forall v. \models [\exists \langle \sigma \rangle. P_{\sigma} \wedge b(\sigma) \wedge v = e(\sigma)] \text{ if } (b) \{C\} [\exists \langle \sigma \rangle. P_{\sigma} \wedge e(\sigma) < v] \quad \forall \sigma. \models [P_{\sigma}] \text{ while } (b) \{C\} [P_{\sigma}] \quad < \text{wf}}{\models [\exists \langle \sigma \rangle. P_{\sigma}] \text{ while } (b) \{C\} [(\exists \langle \sigma \rangle. P_{\sigma}) \wedge \Box (\neg b)]} \text{ (While-}\exists)
\end{array}$$

Fig. 7. Rules from Hyper Hoare Logic [Dardinier and Müller 2023] to reason about while loops. In the rules *WhileSyncTerm* and *While- \exists* , the order $<$ must be *well-founded* (wf). Moreover, $\text{low}(b) \triangleq (\forall \langle \sigma \rangle, \langle \sigma' \rangle. b(\sigma) = b(\sigma'))$ and $\Box(b) \triangleq (\forall \langle \sigma \rangle. b(\sigma))$. Finally, $\models_{\parallel} [P] C [Q]$ corresponds to a *terminating* hyper-triple. Terminating hyper-triples are stronger than normal hyper-triples, in that they additionally require the existence of a terminating execution from any initial state.

red dashed arrow shows the instantiation *forbidden* by our tool, since those instantiations can lead to matching loops. For example, states introduced by \exists^+ -quantifiers can be used to instantiate the \forall^+ -quantifiers of a $\forall^+ \exists^+$ -hyperproperty, which can in turn be used to instantiate the \forall^+ -quantifiers of a \forall^+ -hyperproperty. As a more concrete example, to verify the method *simple* from Fig. 5, the state σ_2 coming from the existential quantifier of the first precondition $\forall \langle \sigma_1 \rangle. \exists \langle \sigma_2 \rangle. \sigma_2(x) > \sigma_1(x)$ can be used to instantiate a \forall -quantifier of the second precondition $\forall \langle \sigma_1 \rangle, \langle \sigma_2 \rangle. \sigma_1(y) = \sigma_2(y)$, since there is a blue arrow from the right of the node $\forall^+ \exists^+$ to the node \forall^+ .

Crucially, states introduced by existential quantifiers that are nested under universal quantifiers *cannot* be used to instantiate the universal quantifiers in $\forall^+ \exists^+$ -hyperproperties (as represented by the red dotted arrow), since this would create a cycle and thus lead to a matching loop, as illustrated above with the first precondition of the method *simple*. As can be seen in Fig. 6, removing this instantiation makes the graph acyclic, which ensures the absence of matching loops.

4 VERIFICATION CONDITIONS FOR LOOPS

In the previous section, we described the verification conditions generated by our verifier for loop-free statements. In this section, we describe how to generate verification conditions for while loops. We first describe, in Sect. 4.1, the different rules offered by Hyper Hoare Logic to reason about while loops, how we can derive verification conditions from them, and how our verifier automatically selects the right rule(s) to apply, based on the context. In Sect. 4.2, we discuss one such particular rule, the *While- $\forall^* \exists^*$* rule, and show that it cannot be used directly for our purpose; a *naive* encoding based on this rule would be unsound. We then present and prove sound (in Isabelle/HOL [Nipkow et al. 2002]) a novel loop rule, more suitable for automated deductive verification, which can be used in the same context. Finally, in Sect. 4.3, we present a technique to automatically frame information around the loop, which overcomes a limitation of these loop rules, and leads to more succinct loop invariants.

4.1 Automatically Generating the Right Verification Conditions

Reasoning about loops in a relational setting is notoriously hard. In the context of deductive verification, our goal is to automatically reason about while loops, while keeping the amount of proof annotations needed from the user to a minimum. As illustrated in Fig. 3, this means that the user should only provide a loop invariant, and optionally a loop variant (**decreases** clause).

Fig. 7 shows the four main rules offered by Hyper Hoare Logic to reason about while loops, where $low(b)$ means that the expression b has the same value in all states (formally $low(b) \triangleq (\forall \langle \sigma \rangle, \langle \sigma' \rangle. b(\sigma) = b(\sigma'))$), $\Box b$ means that the expression b holds in all states (formally $\Box b \triangleq (\forall \langle \sigma \rangle. b(\sigma))$), and $\models_{\parallel} [P] C [Q]$ corresponds to a *terminating* hyper-triple. Terminating hyper-triples are stronger than normal hyper-triples, in that they additionally require the existence of a terminating execution from any initial state. In our tool, we ensure this requirement by proving that all loops in C terminate, through the use of well-founded loop variants.³ As shown by the figure, all four rules use a loop invariant: I in the first three loop rules, and $\exists \langle \sigma \rangle. P_{\sigma}$ in the last rule. Moreover, those rules are non-obvious, which makes it hard for users to know which rule to apply in which context.

In the following, we explain the role of the different rules, how we derive verification conditions from them, and how our verifier automatically chooses the relevant rule(s), based on the user-provided loop invariant and optional loop variant.

Synchronized loop rules. The two first rules, *WhileSync* and *WhileSyncTerm*, apply when all executions exit the loop simultaneously. The key difference between the two rules can be seen in the postconditions of their conclusions: On top of the fact that all states satisfy the negation of the loop guard ($\Box(\neg b)$), the rule *WhileSyncTerm* allows us to assume that the relational invariant I holds after the loop, whereas the rule *WhileSync* allows us to assume only that $I \vee \Box \perp$ holds after the loop, which is weaker than I . The $\Box \perp$ disjunct, which corresponds to the case where the loop does not terminate, is problematic when we want to prove postconditions with top-level existentially-quantified states. In this case, we need to use the rule *WhileSyncTerm*, which requires us to prove that the loop terminates. The latter can be achieved by proving that a well-founded variant e strictly decreases after every iteration.

Verification conditions can be easily derived from these two rules. First, for both rules, we check that the user-provided loop invariant I entails $low(b)$. Then, for the rule *WhileSync*, we separately check the triple $\models [I \wedge \Box b] C [I \vee \Box \perp]$, as described in Sect. 3. For the rule *WhileSyncTerm*, we instantiate e with the user-provided loop variant (required to be an integer expression), and separately check the triple $\models_{\parallel} [I \wedge \Box(b \wedge e = t)] C [I \wedge \Box(0 \leq e < t)]$, where t is a fresh variable. The check $0 \leq e$ ensures that the user-provided variant is well-founded. To ensure that this triple is a terminating hyper-triple, we check that all loops within C are annotated with **decreases** clauses, which ensures termination (provided that verification is successful). Finally, for both rules, we assert that the loop invariant I holds before the loop, and assume that $(I \vee \Box \perp) \wedge \Box(\neg b)$ (rule *WhileSync*) or $I \wedge \Box(\neg b)$ (rule *WhileSyncTerm*) holds after the loop.

Non-synchronized loop rules. The two remaining loop rules from Fig. 7, *While- $\forall^* \exists^*$* and *While- \exists* , can be applied when different executions might exit the loop at different times. In this case, our premises are more complex: We need to reason about the *unrollings* of the while loop, which we achieve by proving a loop invariant over **if** $(b) \{C\}$ (in contrast to C for the synchronized rules). Deriving verification conditions from the rule *While- $\forall^* \exists^*$* is non-trivial, as we explain in Sect. 4.2. For the rule *While- \exists* , assuming that the user-provided loop invariant I is of the form $\exists \langle \sigma \rangle. P_{\sigma}$ (we write P_{σ} to emphasize that this hyper-assertion can mention σ),⁴ and that the user provided a loop

³In theory, we also need to check that **assume** statements do not break this property. By default, our tool leaves this responsibility to the user, since **assume** statements are typically used to restrict non-deterministic assignments to the right range (as done in the example from Fig. 3), which does not break this property. However, our tool provides the more conservative option, disabled by default, to check the absence of **assume** statements within statements whose termination is required.

⁴If I is not of the shape $\exists \langle \sigma \rangle. P_{\sigma}$, and no other rule is applicable, our tool emits an error message to inform the user that the program cannot be verified.

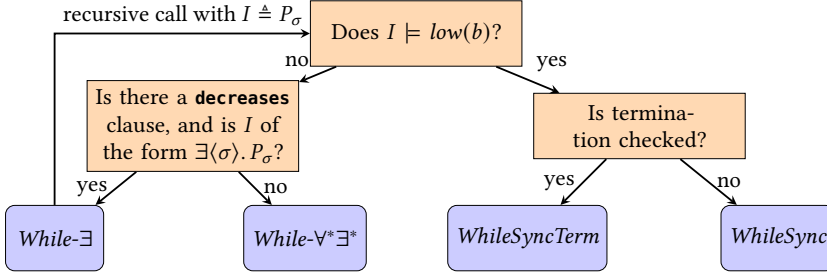


Fig. 8. Automatic loop rule selection to verify a loop **while** (b) $\{C\}$ with the invariant I . To apply the rule *WhileSyncTerm*, the loop **while** (b) $\{C\}$ and all loops nested within C are required to terminate. In contrast, application of the rule *While-∃* does not require termination of nested loops, but only that the outer loop is annotated with a **decreases** clause. Finally, the edge from the rule *While-∃* corresponds to the second premise of this rule, $\forall\sigma. \models [P_\sigma] \text{ while } (b) \{C\} [P_\sigma]$: To check that this premise holds, our tool recursively calls this procedure, which will automatically select a new loop rule adapted to the new loop invariant P_σ .

variant e , we apply the following weakened version of the rule:

$$\frac{\forall v. \models [\exists\langle\sigma\rangle. P_\sigma \wedge b(\sigma) \wedge v = e(\sigma)] \text{ if } (b) \{C\} [\exists\langle\sigma\rangle. P_\sigma \wedge 0 \leq e(\sigma) < v] \quad \forall\sigma. \models [P_\sigma] \text{ while } (b) \{C\} [P_\sigma]}{\models \underbrace{[\exists\langle\sigma\rangle. P_\sigma]}_I \text{ while } (b) \{C\} \underbrace{[(\exists\langle\sigma\rangle. P_\sigma) \wedge \Box(\neg b)]}_I}$$

As before, this version specializes the well-founded order $<$ to be the canonical well-founded order over natural numbers. Note that, in both premises, v and σ are *meta-variables*, i.e., there is not one value of v (in the first premise) or σ (in the second premise) per state, but rather there is one per *set* of states.

In practice, to check the first premise, we use a fresh unconstrained variable v , and assume that the set of states and the variable v together satisfy the precondition $\exists\langle\sigma\rangle. P_\sigma \wedge b(\sigma) \wedge v = e(\sigma)$, and check (after the encoding of **if** (b) $\{C\}$) that the set of states and the variable v together satisfy the postcondition $\exists\langle\sigma\rangle. P_\sigma \wedge 0 \leq e(\sigma) < v$. Checking the second premise is more complicated, since it requires to reason about the same **while** loop. However, note that the precondition (and postcondition) of this premise, P_σ , is smaller than the precondition (and postcondition) of the conclusion of the rule, $\exists\langle\sigma\rangle. P_\sigma$. Our tool automatically generates the verification conditions for this premise, using the approach described in this section, by automatically selecting the right loop rule based on the new loop invariant P_σ and the same loop variant e .

Automatically selecting the right loop rule(s). As explained at the start of this section, using only the user-provided loop invariant I and optional loop variant e , our tool automatically selects the right loop rule(s) to apply, as depicted in Fig. 8. First, we check whether the loop invariant guarantees that all executions will exit the loop simultaneously, by whether $I \models \text{low}(b)$ holds. If so, we apply one of the two synchronized loop rules, *WhileSync* or *WhileSyncTerm*, depending on whether the user provided a loop variant for this loop and all loops nested within. The reason is that these rules, when applicable, are always better than the non-synchronized loop rules. The rule *WhileSyncTerm* is the most powerful (when it applies), because its premise only requires to reason about C , which is easier than reasoning about **if** (b) $\{C\}$, and the postcondition of its conclusion, $I \wedge \Box(\neg b)$, is not weaker than the postcondition given by any other rule (Sect. 4.2 will make clearer why the postcondition of the rule *While-∀*∃** is weaker). When termination is not checked (i.e., no

```

687 method naive_encoding(t: Int, n: Int) returns (x: Int)
688   requires  $\exists \langle \sigma_1 \rangle. \sigma_1(t) = 1$ 
689   requires  $\forall v. v \geq 0 \Rightarrow \exists \langle \sigma_2 \rangle. \sigma_2(t) = 2 \wedge \sigma_2(n) = v$ 
690   ensures  $\exists v. \forall \langle \sigma \rangle. \sigma(x) \leq v$  // this postcondition does not hold
691 {
692   x := 0
693   while (t = 1 || x < n)
694     invariant  $\exists \langle \sigma_1 \rangle. \sigma_1(t) = 1$ 
695     invariant  $\forall \langle \sigma_1 \rangle, \langle \sigma_2 \rangle. \sigma_1(t) = 1 \Rightarrow \sigma_1(x) \geq \sigma_2(x)$ 
696   {
697     x := x + 1
698   }
699 }

```

Fig. 9. An example showing why a *naive* encoding based on the rule $\text{While-}\forall^*\exists^*$ would be unsound.

loop variant is provided), the choice is only between the rules *WhileSync* and *While- $\forall^*\exists^*$* , and our tool applies *WhileSync* whenever possible, for similar reasons.

When $I \not\models \text{low}(b)$, we apply one of the two non-synchronized loop rules, *While- $\forall^*\exists^*$* or *While- \exists* , depending on the shape of the loop invariant I . When I is of the form $\forall^+\exists^*$, we can only apply the rule *While- $\forall^*\exists^*$* , because our invariant is not of the form $\exists \langle \sigma \rangle. P$, required by the rule *While- \exists* . Similarly, when I is of the form $\exists^+\forall^+$, we can only apply the rule *While- \exists* , because I does not satisfy the syntactic restriction from the rule *While- $\forall^*\exists^*$* . Thus, there exists a *choice* between those two loop rules only when I has the shape \exists^+ and a loop variant is provided (otherwise the rule *While- \exists* cannot be applied). In this case, the rule *While- \exists* is more powerful, since it easily allows proving that the existentially-quantified states in I will still exist after the while loop, thanks to the loop variant. As a concrete example, let $I \triangleq \exists \langle \sigma \rangle. \sigma(x) = \sigma(y)$. We can use the rule *While- \exists* with $P_\sigma \triangleq (\sigma(x) = \sigma(y))$, which gives us the desired postcondition $\exists \langle \sigma \rangle. \sigma(x) = \sigma(y)$ in its conclusion. In contrast, the postcondition of the conclusion of the rule *While- $\forall^*\exists^*$* is some hyper-assertion Q , such that $\models [\exists \langle \sigma \rangle. \sigma(x) = \sigma(y)] \text{ assume } \neg b [Q]$ holds. In particular, we can get our desired postcondition $\exists \langle \sigma \rangle. \sigma(x) = \sigma(y)$ only if $\sigma(x) = \sigma(y)$ implies $\neg b$, which will typically not be the case (because $\exists \langle \sigma \rangle. \sigma(x) = \sigma(y)$ is our loop invariant, which has to already hold *before* the loop). Thus, when applicable, our tool applies the rule *While- \exists* over the rule *While- $\forall^*\exists^*$* , which then recursively applies the same automatic loop rule selection with the smaller loop invariant P , as shown in Fig. 8.

4.2 $\forall^*\exists^*$ -Hyperproperties

In many cases, the only loop rule that can be applied is the rule *While- $\forall^*\exists^*$* . then our tool tries to apply the rule *While- $\forall^*\exists^*$* from Fig. 7. At first glance, this rule looks straightforward to automate: The premise $\models [I] \text{ if } (b) \{C\} [I]$ can be checked separately using the user-provided loop invariant I , and the postcondition Q can be obtained by considering a fresh set of states after the loop, assuming that it satisfies I , and then encoding *assume $\neg b$* . This approach allows us to derive Q *semantically*. The key difficulty then lies in checking the *syntactic* restriction (no universal state-quantifier should occur under an existential quantifier) on Q . One naive idea would be to enforce this syntactic restriction for the loop invariant I , and hope that the rule is still sound. Unfortunately, this surprisingly results in an unsound encoding, as we illustrate next.

Unsoundness of the naive encoding. Consider the method `naive_encoding` from Fig. 9. Depending on the value of t , this method will either loop forever (if $t = 1$) or simply increment x until $x = n$ (if $n \geq 0$ and $t = 2$). Our precondition⁵ requires the existence of a state that will loop forever ($t = 1$), and, for all possible non-negative values v of n , the existence of a state that will do n iterations until $x = n$. Thus, after the loop, we should have a set of states that contains at least one state for each possible non-negative value v of x . In particular, we should *not* be able to prove the postcondition. However, using the encoding described above, we *are able* to prove this incorrect postcondition!

Indeed, the first premise of the rule $\text{While-}\forall^*\exists^*$, $\models [I] \text{ if } (t = 1 \vee x < n) \{x := x + 1\} [I]$, holds, since any state σ_1 with $t = 1$ will enter the if-branch, and thus σ_1 will keep having the maximal value (among all executions) for x . Moreover, the loop invariant I satisfies the syntactic restriction (no $\forall(_)$ appears under any existential quantifier), and I clearly holds before the loop, since $x = 0$ in all states. Finally, let us consider what happens after the loop, and why this encoding allows us to derive the wrong postcondition. Let S be a set of states that satisfies the loop invariant I , and let S' be the set of states obtained by executing `assume` $\neg(t = 1 \vee x < n)$ in all states from S . Note that S' corresponds to the subset of states from S that satisfy $t \neq 1 \wedge x \geq n$. From I , we learn that there is a state σ_1 in S where $t = 1$, and that this state σ_1 has the maximum value for x among all states in S . Thus, there exists an upper bound v for the value of x in all states from S , namely $v \triangleq \sigma_1(x)$. Since S' is a subset of S , this upper bound v is also an upper bound for the value of x in all states from S' , which corresponds to the incorrect postcondition.

A new rule for automating \forall^\exists^* -hyperproperties.* The previous example shows that deriving the postcondition Q *semantically* from I , while checking the *syntactic* restriction on the loop invariant I , is unsound. We solve this issue by deriving the postcondition Q *syntactically* from I while enforcing the *syntactic* restriction on I . This allows us to obtain a sound rule, which can be automated in a straightforward way. We obtain the postcondition from I , which we write $\Theta_{\neg b}(I)$, by recursively replacing all instances of $\exists\langle\sigma\rangle.P$ with $\exists\sigma.P \wedge (\neg b \Rightarrow \langle\sigma\rangle)$. That is, the postcondition ensures that the existentially-quantified states in I *exist*, but they are not guaranteed to belong to the set of states after the loop: They belong to the set of states after the loop if they satisfy the negation $\neg b$ of the loop guard. In the example from Fig. 9, we obtain from I the postcondition $\Theta_{\neg b}(I) = (\exists\sigma_1.\sigma_1(t) = 1 \wedge (\neg(\sigma_1(t) = 1 \vee \sigma_1(x) < \sigma_1(n)) \Rightarrow \langle\sigma_1\rangle)) \wedge (\forall\langle\sigma_1\rangle, \langle\sigma_2\rangle.\sigma_1(t) = 1 \Rightarrow \sigma_1(x) \geq \sigma_2(x))$, which does not entail the wrong postcondition anymore (recall that $\langle\sigma_1\rangle$ is syntactic sugar for $\lambda S.\sigma_1 \in S$). Indeed, while we still learn the existence of a state σ_1 where $t = 1$, we do not learn that σ_1 belongs to the set of states after the loop, because we cannot prove $\neg(\sigma_1(t) = 1 \vee \sigma_1(x) < \sigma_1(n))$, and, thus, we cannot conclude that $\forall\langle\sigma_2\rangle.\sigma_1(x) \geq \sigma_2(x)$.

We have proven in Isabelle/HOL [Nipkow et al. 2002] that this novel rule, which our tool leverages, is sound:

THEOREM 1. Novel loop rule for $\forall^*\exists^*$ -hyperproperties. *Let C be a program statement, b a program expression, and I a (syntactic) hyper-assertion such that I contains no $\forall(_)$ after any \exists . If $\models [I] \text{ if } (b) \{C\} [I]$ holds, then $\models [I] \text{ while } (b) \{C\} [\Theta_{\neg b}(I) \wedge \Box(\neg b)]$ holds.*

PROOF SKETCH. To prove this result, we use the fact that $\text{sem}(\text{while } (b) \{C\}, S)$, the semantics of the while loop given a set of initial states S , can be seen as the limit of $\text{sem}([\text{if } (b) \{C\}]^n; \text{assume } \neg b, S)$ as n goes to infinity, where $[\text{if } (b) \{C\}]^n$ represents the statement `if` $(b) \{C\}$ sequentially composed with itself n times. More formally:

$$\text{sem}(\text{while } (b) \{C\}, S) = \bigcup_{n \in \mathbb{N}} \text{sem}([\text{if } (b) \{C\}]^n; \text{assume } \neg b, S) \quad (*)$$

⁵Note that no precondition is actually required for the naive encoding to be unsound on this example, but we use one to simplify the explanations.

method framing1(x : Int) returns (y : Int)	method framing2(x : Int) returns (y : Int)
requires $\forall \langle \sigma_1 \rangle, \langle \sigma_2 \rangle. \sigma_1(x) = \sigma_2(x)$	requires $\exists \langle \sigma \rangle. \forall \langle \sigma' \rangle. \sigma(x) \geq \sigma'(x)$
requires $\forall \langle \sigma \rangle. \sigma(x) \geq 0$	requires $\forall \langle \sigma \rangle. \sigma(x) \geq 0$
ensures $\forall \langle \sigma_1 \rangle, \langle \sigma_2 \rangle. \sigma_1(y) = \sigma_2(y)$	ensures $\exists \langle \sigma \rangle. \forall \langle \sigma' \rangle. \sigma(y) \geq \sigma'(y)$
{	{
$y := 0$	$y := 0$
while ($y < x$)	while ($y < x$)
invariant $\forall \langle \sigma \rangle. \sigma(y) \leq \sigma(x)$	invariant $\forall \langle \sigma \rangle. \sigma(y) \leq \sigma(x)$
decreases $x - y$	decreases $x - y$
{ $y := y + 1$ }	{ $y := y + 1$ }
}	}

Fig. 10. An example from HYPR that requires automatic framing to be successfully verified.

In other words, every state after the loop must have exited the loop after n iterations, for some n . In particular, note that the sequence of sets $(\text{sem}([\text{if } (b) \{C\}]^n; \text{assume } \neg b, S))_{n \in \mathbb{N}}$ is non-decreasing. That is, the set of states that exit the loop in the first n iterations can only grow when n grows.

We then prove the following property $\mathcal{P}(I)$, by structural induction over the syntactic hyperassertion I : "For any non-decreasing sequence $(S_n)_{n \in \mathbb{N}}$ of set of states, if I contains no $\forall(_)$ after any \exists , and if $\forall n. S_n \models \Theta_{-b}(I)$, then $(\bigcup_n S_n) \models \Theta_{-b}(I)$ holds". The theorem follows from this property and the aforementioned identity (*). In the following, we discuss four cases of the induction; all other cases are trivial.

The cases for $\mathcal{P}(\exists y. I)$ and $\mathcal{P}(\exists \langle \sigma \rangle. I)$ are straightforward: By the syntactic restriction, we know that I contains no $\forall(_)$, and so does $\Theta_{-b}(I)$. Intuitively, this means that $\Theta_{-b}(I)$ only cares about the *existence* of states, and thus $\Theta_{-b}(I)$ grows monotonically (which can be proven by an additional trivial induction on I): If it is satisfied by a set of states, then it will be satisfied by any superset of this set. Since it is satisfied by all S_n , it is also satisfied by their union.

For the case $\mathcal{P}(\forall \langle \sigma \rangle. I)$, we get to assume (1) $\mathcal{P}(I)$ and (2) $\forall n. S_n \models \forall \langle \sigma \rangle. \Theta_{-b}(I)$, and we want to prove $(\bigcup_n S_n) \models \forall \langle \sigma \rangle. \Theta_{-b}(I)$. To prove this, let σ be a state in $\bigcup_n S_n$, and let us prove that $(\bigcup_n S_n), \sigma \models \Theta_{-b}(I)$ (which informally means that the previously existentially-quantified state σ is instantiated in $\Theta_{-b}(I)$ to the concrete state)⁶. Because $\sigma \in \bigcup_n S_n$, there exists a k such that $\sigma \in S_k$. Let S' such that $\forall n. S'_n = S_{n+k}$. Because $\forall n. S'_n, \sigma \models \Theta_{-b}(I)$, we can use the induction hypothesis $\mathcal{P}(I)$ to get that $(\bigcup_n S'_n) \models \Theta_{-b}(I)$. Finally, notice that $(\bigcup_n S_n) = (\bigcup_n S'_n)$, because S is non-decreasing, which concludes the case.

For the case $\mathcal{P}(I_1 \vee I_2)$, we get to assume (1) $\mathcal{P}(I_1)$, (2) $\mathcal{P}(I_2)$, and (3) $\forall n. S_n \models \Theta_{-b}(I_1) \vee \Theta_{-b}(I_2)$, and we want to prove $(\bigcup_n S_n) \models \Theta_{-b}(I_1) \vee \Theta_{-b}(I_2)$. By (3), we know that $\Theta_{-b}(I_1) \vee \Theta_{-b}(I_2)$ is true infinitely often, which implies that either $\Theta_{-b}(I_1)$ is true infinitely often, or $\Theta_{-b}(I_2)$ is true infinitely often. Without loss of generality, let us assume that $\Theta_{-b}(I_1)$ is true infinitely often, and let S' be an infinite subsequence of S such that $\forall n. S'_n \models \Theta_{-b}(I_1)$. By the induction hypothesis $\mathcal{P}(I_1)$, we get that $(\bigcup_n S'_n) \models \Theta_{-b}(I_1)$. Moreover, because S is non-decreasing, we get that $(\bigcup_n S_n) = (\bigcup_n S'_n)$, which concludes the case. \square

4.3 Automatic Framing

In the previous subsections, we have shown how we derived verification conditions from the loop rules offered by Hyper Hoare Logic. However, using those loop rules on their own (and not in conjunction with other rules as we show below) has the limitation that only the information

⁶In our mechanization, we use de Bruijn indices to handle quantifiers, which we ignore in this paper for simplicity.

provided by the loop invariant is preserved, as we illustrate with the examples in Fig. 10. Consider the method `framing1` on the left of the figure, which increments y in a loop until $x = y$. We want to prove that if x has the same initial value in all executions, then y will have the same final value in all executions. Using the standard (unary) loop invariant $\forall \langle \sigma \rangle. \sigma(y) \leq \sigma(x)$, we can easily prove that, after the loop, $x = y$ in all states (1). Moreover, since all executions have the same value for x before the loop, and since x is not modified by the loop, all executions will still have the same value for x after the loop (2). By conjoining (1) and (2), we get the postcondition.

Unfortunately, the loop encodings presented so far are only able to prove (1), but not (2), since they only assume (a property derived from) the loop invariant after the loop. Because our loop invariant does not mention that x has the same value in all executions, this piece of information is lost after the loop.

One way to solve this particular problem is to add this information to the loop invariant, by conjoining $\forall \langle \sigma_1 \rangle, \langle \sigma_2 \rangle. \sigma_1(x) = \sigma_2(x)$ to it. This solution is cumbersome for the user, who is required to write longer invariants, by adding information not relevant for the loop (but only relevant for the postcondition later). Another way to solve this issue is to use the following rule from Hyper Hoare Logic (where $\text{mod}(C)$ represents the variables modified by C and $\text{fv}(F)$ the (program) variables that appear in F), which allows propagating information about variables not modified by the loop after the loop:

$$\frac{\models [P] \ C \ [Q] \quad \text{no } \exists \langle _ \rangle \text{ in } F \quad \text{mod}(C) \cap \text{fv}(F) = \emptyset}{\models [P \wedge F] \ C \ [Q \wedge F]} \text{ (FrameSafe)}$$

Since x is not modified by the loop, we can use this rule with $F \triangleq (\forall \langle \sigma_1 \rangle, \langle \sigma_2 \rangle. \sigma_1(x) = \sigma_2(x))$ to solve our issue. Our goal is to use this rule to *automatically* frame information around the loop, without requiring the user to provide F .

We achieve this by recording the set of states before the loop in an auxiliary variable S_p , and by adding (after the loop) the assumption that, for each state in the set of states after the loop, there must exist a state in S_p with the same values for all variables not modified by C :

```

 $S_p := S$ 
... // loop encoding
assume  $\forall \sigma' \in S. \exists \sigma \in S_p. (\forall x. x \notin \text{fv}(C) \Rightarrow \sigma(x) = \sigma'(x))$  // (F-0X)

```

Intuitively, adding this assumption is sound because every state σ' after the loop corresponds to the final state of an execution of C in an initial state σ from S_p , and thus σ and σ' must have the same values for the variables not modified by C . Formally, this encoding is justified by the following straightforward lemma:

LEMMA 1. $\forall \sigma' \in \text{sem}(C, S). \exists \sigma \in S. \forall x \notin \text{mod}(C). \sigma(x) = \sigma'(x)$

This encoding is stronger than *any* possible application of the rule *FrameSafe*, since the former logically implies the latter (for any frame F). To see why it solves the issue from our example, consider two states σ'_1 and σ'_2 that belong to the set of states S after the loop. From the assumption (F-0X), we get the existence of two states σ_1 and σ_2 from S_p , such that $\sigma_1(x) = \sigma'_1(x)$ and $\sigma_2(x) = \sigma'_2(x)$. Because of the precondition, we know that $\sigma_1(x) = \sigma_2(x)$, and thus can conclude that $\sigma'_1(x) = \sigma'_2(x)$.

Framing hyperproperties with existentially-quantified states. Note that the rule *FrameSafe* has the restriction that F is not allowed to existentially quantify over states. To see why this is a limitation, consider as an example the method `framing2` on the right of Fig. 10. This method has the same body and loop invariant as method `framing1`, but we now want to prove that if there is an execution

whose initial value x is maximal (among all executions), then there should exist an execution whose final value for y is also maximal. In this case, we would like to apply the rule *FrameSafe* with the frame $F \triangleq (\exists \langle \sigma \rangle. \forall \langle \sigma' \rangle. \sigma(x) \geq \sigma'(x))$, but cannot because of this restriction.

To overcome this limitation, Hyper Hoare Logic provides the following rule, which lifts this restriction:

$$\frac{\text{mod}(C) \cap \text{fv}(F) = \emptyset \quad \models_{\Downarrow} [P] \ C \ [Q] \quad F \text{ is a syntactic hyper-assertion}}{\models_{\Downarrow} [P \wedge F] \ C \ [Q \wedge F]} \text{ (Frame)}$$

This rule requires however to prove a stronger triple, the *terminating* hyper-triple $\models_{\Downarrow} [P] \ C \ [Q]$, which must ensure the existence of a terminating execution from any initial state. In our tool, we can ensure that a triple around a loop **while** (b) $\{C\}$ is terminating as long as C contains no **assume** statements, and this loop and all nested loops in C terminate, i.e., have been annotated with a **decreases** clause. This is for example the case for the loop in method `framing2`. When those two conditions hold, it is sound to strengthen the previous encoding with the additional assumption (F-UX), as follows:

```

 $S_p := S$ 
... // loop encoding
assume  $\forall \sigma' \in S. \exists \sigma \in S_p. (\forall x. x \notin \text{fv}(C) \Rightarrow \sigma(x) = \sigma'(x))$  // (F-0X)
assume  $\forall \sigma \in S_p. \exists \sigma \in S. (\forall x. x \notin \text{fv}(C) \Rightarrow \sigma(x) = \sigma'(x))$  // (F-UX)

```

Together, those two assumptions are stronger than the application of the rule *Frame* for any frame F . For example, emitting those two assumptions together lets us automatically derive that $\exists \langle \sigma \rangle. \forall \langle \sigma' \rangle. \sigma(x) \geq \sigma'(x)$ holds after the loop in our example. However, we have noticed in practice that emitting the assumption (F-UX) does not interact well with the encoding described in Sect. 3.2, and might result in matching loops. Thus, our tool provides an option to emit this second assumption (when applicable), which is disabled by default.

5 IMPLEMENTATION AND EVALUATION

We implemented HYPR, a deductive program verifier for hyperproperties, on top of Viper; that is, HYPR takes as input a text file, translates it into a Viper program (as described in Sect. 3), calls the Viper verifier to verify this program, and then translates the output (successful verification or error messages) back to the user.⁷

We evaluated HYPR on a diverse set of examples, which includes many examples from the literature, to answer the following questions:

(RQ1) Can our tool (dis-)prove hyperproperties of different types, namely \forall^* , \exists^* , $\forall^* \exists^*$, and $\exists^* \forall^*$?

(RQ2) How many lines of proof annotations are needed by our tool?

(RQ3) Can our tool verify complex examples in a reasonable amount of time?

In summary, our evaluation shows that our tool can efficiently (dis-)prove hyperproperties of different types with a reasonable amount of proof annotations, and it can do so within a reasonable amount of time. In the following, we describe how we selected our benchmarks, how we ran the experiments, and present and discuss the results.

Benchmarks. To evaluate HYPR, we used the benchmarks from DESCARTES [Sousa and Dillig 2016], HyPA [Beutner and Finkbeiner 2022], ORHLE [Dickerson et al. 2022], and PCSAT [Unno et al. 2021]: We selected a subset of their publicly-available benchmarks and translated them into the

⁷Viper actually provides two verifiers, one based on symbolic execution, and one based on BOOGIE [Leino 2008]. Our tool uses the two Viper verifiers to verify the generated Viper program, and reports the first successful verification result.

Type of hyperproperty	Source	Files		Verification time		Annotations
		no.	Mean (LoC)	Mean (s)	Median (s)	
\forall^*	Descartes	15	129	2.74	1.93	0
	PCSat	4	25	2.80	1.72	3
	Overall	19	107	2.75	1.89	1
\exists^*	Descartes	8	81	15.13	6.84	0
	ORHLE	6	29	9.78	1.39	8
	Overall	14	59	12.83	4.04	4
$\forall^*\exists^*$	ORHLE	28	21	1.80	1.21	2
	HyPa	8	14	1.31	1.19	3
	PCSat	1	22	1.31	1.31	2
	Overall	37	19	1.68	1.20	2
$\exists^*\forall^*$	ORHLE	15	25	2.85	1.19	2

Fig. 11. Results of our evaluation.

programming language supported by our tool to form our benchmarks. We selected the benchmarks based on the following criteria:

- (1) For DESCARTES, ORHLE, HyPA, and PCSAT, we ignored the benchmarks that use data structures not supported by our tool, such as arrays.
- (2) For DESCARTES, we ignored the benchmarks that use objects with more than 3 fields, since translating fields into the language supported by our tool is cumbersome.
- (3) For ORHLE, HyPA, and PCSAT, we ignored the benchmarks that prove hyperproperties over *different* programs, since our tool does not support this.
- (4) For PCSAT, we only selected the benchmarks that do not require reasoning about co-termination, since our tool does not support this.

For each selected benchmark, we translated it to the syntax accepted by HYPRA. To obtain hyper-triples semantically equivalent to the original specifications, we used the formal translations given by Dardinier and Müller [2023]. In addition, we annotated the translated benchmarks with loop variants, loop invariants and hints when necessary.

For invalid benchmarks that fail to prove \forall^* or $\forall^*\exists^*$ hyperproperties, we also formally disprove them. To do so, we strengthened the preconditions and proved the negation of the original postconditions [Dardinier and Müller 2023, Theorem 4]. In particular, this allows us to obtain benchmarks that prove $\exists^*\forall^*$ -hyperproperties, since the benchmarks from those tools do not contain any.

In total, we obtained 85 benchmarks. Fig. 11 provides more details about the selected benchmarks.

Experimental Setup. We ran HYPRA to verify the translated benchmarks on a MacBook Pro running macOS Ventura 13.3 with a 2.3 GHz 8-Core Intel Core i9 processor and 32 GB RAM. Each benchmark was run with 3 repetitions. For each run, we recorded the verification result and runtime. In the end, we checked that the verification results in all runs were consistent, and also computed the average verification time for each benchmark.

Results. The results of our evaluation are shown in Fig. 11. As we can see, HYPRA can handle not only all \forall^* , \exists^* and $\forall^*\exists^*$ -hyperproperties that other verifiers can handle, but also $\exists^*\forall^*$ -hyperproperties, which no other existing verifier supports.

Although verification using HYPRA is not fully automated, it only requires a reasonable amount of proof annotations from users, which is evidenced by the last column of Fig. 11.

Moreover, HYBRA is quite efficient in general. On average, it took HYBRA 338 seconds to run the entire benchmark suite composed of 85 programs. For 87% of the benchmarks, verification finished within 5 seconds. In some rare cases, the runtime was relatively long, with the maximum runtime around 45 seconds. This is not unexpected, since some of those benchmarks have very complex commands (such as lots of nested conditional statements) and specifications (such as preconditions and postconditions of the shape $\exists\exists\exists\forall\forall\forall$).

In summary, our evaluation demonstrates that HYBRA can effectively verify hyperproperties of different types with a reasonable amount of proof annotations and within a reasonable amount of time.

6 RELATED WORK

In this section, we first cover related program logics for hyperproperties (Sect. 6.1), and then tools and approaches for automatically verifying hyperproperties (Sect. 6.2).

6.1 Program Logics for Hyperproperties

As discussed in Sect. 1, many logics to prove that a program satisfies a safety hyperproperty [Clarkson and Schneider 2008] have been proposed over the years [Naumann and Ngo 2019]. Many of those logics actually prove *relational* properties, i.e., properties that relate the executions of several (potentially different) programs. For example, Relational Hoare Logic (RHL) [Benton 2004] extends Hoare Logic [Floyd 1967; Hoare 1969] to reason about $\forall\forall$ -properties relating the executions of two programs, e.g., to prove correct some program transformations. RHL has later been extended to handle more complex programs. For example, RHL has been combined with separation logic [Reynolds 2002], to support relational properties between two heap-manipulating programs in a modular way [Yang 2007]. RHL has also been extended to reason about higher-order programs [Aguirre et al. 2017]. Several program logics [Amtoft et al. 2006; Costanzo and Shao 2014; Eilers et al. 2023; Ernst and Murray 2019] have been designed specifically to prove non-interference [Volpano et al. 1996] properties, a particular case of 2-safety hyperproperties, and thus support properties for all pairs of executions of a *single* program.

Several important safety hyperproperties are not 2-safety hyperproperties, but k -safety hyperproperties for $k > 2$, such as transitivity ($k = 3$) or associativity ($k = 4$). Sousa and Dillig [2016] have proposed Cartesian Hoare Logic (CHL) to reason about k -safety properties for any fixed k . D’Ousualdo et al. [2022] have identified several limitations of CHL when trying to compose together proofs of different k -safety properties, and have proposed a novel weakest-precondition calculus to overcome these limitations.

All aforementioned logics are *overapproximate* logics, that is, they work by overapproximating the set of executions, which is sufficient for proving safety hyperproperties. However, hyperproperties outside the safety class require proving the *existence* of relevant executions, which requires *underapproximation*. Underapproximate logics include Reverse Hoare Logic [de Vries and Koutavas 2011] and Incorrectness Logic [O’Hearn 2019], which are useful to prove reachability properties or the existence of bugs. Underapproximate logics have proven useful to justify the formal foundations of industrial bug-finding tools [Blackshear et al. 2018; Distefano et al. 2019; Gorogiannis et al. 2019; Le et al. 2022]. Several logics combine over- and underapproximate reasoning for *single* executions, such as Dynamic Logic [Harel 1979], Outcome Logic [Zilberstein et al. 2023], Exact Separation Logic [Maksimović et al. 2023], or Local Completeness Logic [Bruni et al. 2021]. Recently, Murray [2020] has proposed a program logic for $\exists\exists$ -hyperproperties, to prove the presence of insecurity in a program.

Finally, several program logics combining over- and underapproximate reasoning for relating *multiple* executions have been proposed, to reason about $\forall^*\exists^*$ or $\exists^*\forall$ -hyperproperties. Maillard et al.

[2019] present a general framework for defining relational program logics (for two executions of two potentially different programs), which can be instantiated for $\forall\exists$ -properties. RHLE [Dickerson et al. 2022] combines an underapproximate and an overapproximate Hoare logic, to support relational $\forall^*\exists^*$ -properties between (potentially different) programs. Antonopoulos et al. [2023] present BiKAT, an extension of KAT [Kozen 1997], useful to reason about alignment in the context of relational verification, and derive from BiKAT inference rules for $\forall\forall$ and $\forall\exists$ -properties. The authors also show that BiKAT can in principle also be used for $\exists\exists$ and $\exists\forall$ -properties. Finally, Dardinier and Müller [2023] present Hyper Hoare Logic (HHL), on which we base our approach. Unlike the previously-mentioned logics, HHL is tailored to hyperproperties (i.e., properties relating the executions of a *single* program), and does not support relational properties between different programs. However, HHL supports a larger class of hyperproperties than these logics, since it supports hyperproperties with arbitrary quantifier alternation, capturing both $\forall^*\exists^*$ and $\exists^*\forall^*$ -hyperproperties.

6.2 Automated Verification of Hyperproperties

Deductive Verification. Deductive program verifiers (or deductive verifiers for short) are tools that, given as input a program, a specification, and proof hints (such as loop invariants), try to automatically construct a proof in a given program logic that the program satisfies the specification. Many deductive verifiers based on SMT solvers (such as Z3 [de Moura and Bjørner 2008]) have been developed for verifying *safety properties*, i.e., properties that should hold for all *individual* executions, such as Boogie [Leino 2008], Why3 [Filliâtre and Paskevich 2013], DAFNY [Leino 2010] (based on dynamic frames [Kassios 2006]), or VIPER [Müller et al. 2016] (based on separation logic [Reynolds 2002]).

The problem of verifying that a program satisfies a k -safety hyperproperty can be reduced to the problem of verifying that a product program [Barthe et al. 2011; Terauchi and Aiken 2005] satisfies a safety property, where the product program is for example obtained by composing sequentially k renamed copies of the original program. The product program can then be verified using deductive verifiers tailored for safety properties. Eilers et al. [2019] show how to treat method calls modularly in this context, allowing methods to have relational preconditions and postconditions, similar to the \forall^* -specifications shown in Sect. 2 (for example in Fig. 1).

Deductive verifiers specifically targeting hyperproperties have been developed as well. Those include WHYREL [Nagasamudram et al. 2023], SecC (based on SecCSL) [Ernst and Murray 2019], and HYPERVIPER (based on CommCSL) [Eilers et al. 2023] for non-interference [Volpano et al. 1996] (a 2-safety hyperproperty), DESCARTES (based on Cartesian Hoare Logic) [Sousa and Dillig 2016] for k -safety hyperproperties, and ORHLE (based on RHLE) [Dickerson et al. 2022] for $\forall^*\exists^*$ -hyperproperties. As our evaluation shows, our tool handles well the benchmarks from DESCARTES and ORHLE, and can even disprove invalid ones. Compared to ORHLE, the closest to our work, our tool HYPRA is more expressive, since it also supports for example $\exists^*\forall^*$ -hyperproperties, and supports reasoning about runtime errors. Our tool is also more flexible, since it allows the user to write explicit quantifiers in the assertion language itself, and thus allows to combine different types of hyperproperties in the same proof, whereas ORHLE requires the user to fix the quantification scheme in advance. Moreover, even for $\forall^*\exists^*$ -hyperproperties, our tool supports reasoning about more complex proof patterns, such as while loops where different executions might exit at different iterations.

Other approaches have been developed to automatically verify hyperproperties [Assaf et al. 2017; Barthe et al. 2019; Farzan and Vandikas 2019; Itzhaky et al. 2024; Unno et al. 2021]. For example, Assaf et al. [2017] use abstract interpretation [Cousot and Cousot 1977] to verify different hypersafety properties related to information flow, including some safety hyperproperties that are not k -safety for any k . To achieve this, they present a hypercollecting semantics, similar in

spirit to the function *sem* (Def. 1) from Hyper Hoare Logic. Unno et al. [2021] present PCSAT, a tool based on a generalization of Constrained Horn Clauses [Bjørner et al. 2015] to automatically verify k -safety hyperproperties, and more complex hyperproperties such as termination-sensitive non-interference [Volpano and Smith 1997], and generalized non-interference [McCullough 1987; McLean 1996]. As shown in our evaluation, our tool HYPERA can handle all the benchmarks from PCSAT that fall in our supported subset of programs, with a reasonable amount of proof annotation and in reasonable time. Extending HYPERA to reason about properties such as termination-sensitive non-interference is future work.

Finally, *temporal* logics to express hyperproperties have been proposed, such as HyperLTL and HyperCTL* [Clarkson et al. 2014], and model checking [Clarke 1997] algorithms to check whether finite-state systems satisfy hyperproperties expressed in these temporal logics have been proposed [Finkbeiner et al. 2015]. For example, Hsu et al. [2021] have proposed algorithms for *bounded* model checking, Coenen et al. [2019] proposed model checking algorithms for $\forall^*\exists^*$ -hyperproperties, and Beutner and Finkbeiner [2023] proposed an explicit-state model checking algorithm that is complete for HyperLTL and for hyperproperties with arbitrary quantifier alternations. Beutner and Finkbeiner [2022] have also shown that model checking techniques for $\forall^*\exists^*$ can be applied to infinite-state systems, by using predicate abstraction.

7 CONCLUSION AND FUTURE WORK

In this paper, we have presented a novel approach for the deductive verification of hyperproperties, including hyperproperties of the shape $\forall^*\exists^*$ and $\exists^*\forall^*$. Our approach, based on an extension of Hyper Hoare Logic [Dardinier and Müller 2023] to reason about runtime errors (Sect. 2.2) and implemented in the tool HYPERA, tracks an underapproximation and an overapproximation of the set of reachable states and the set of error states (Sect. 3.1), and combines those approximations with some carefully-designed restrictions to avoid matching loops (Sect. 3.2). Moreover, our tool is able to generate automatically verification conditions for loops, based on a user-provided loop invariant and optional variant. To achieve this, our tool automatically selects the right loop rule to apply based on the context (Sect. 4.1), which includes a novel loop rule for proving $\forall^*\exists^*$ -hyperproperties (Sect. 4.2), proved sound in Isabelle/HOL. Finally, our tool is able to automatically frame hyperproperties untouched by loops (Sect. 4.3), leading to more concise loop invariants. Our evaluation (Sect. 5) shows that HYPERA can prove a large class of hyperproperties for a large class of programs, in a reasonable amount of time and with a reasonable amount of proof annotation.

Our work shows that it is possible to build an effective automated deductive program verifier for a large class of hyperproperties, and it opens several avenues for future work. Currently, our tool does not support heap-manipulating programs. One interesting research direction would thus be to extend Hyper Hoare Logic to support reasoning about heap-manipulating programs in a modular way (for example by borrowing concepts from separation logic [Reynolds 2002]), and then extend our approach based on this extended logic. Another interesting research direction would be to extend our approach to support relational properties with arbitrary quantifier alternation between *different* programs, which could for example allow to prove that a program *does not* refine another one, an $\exists\forall$ -property. Finally, similar to how we extended Hyper Hoare Logic to support runtime errors, it would be interesting to explore an extension of Hyper Hoare Logic to reason about termination and non-termination *in the assertion language itself*. On top of proving termination (which our tool is already capable of), this would enable proving non-termination, and hyperproperties such as co-termination (e.g., to prove that observing non-termination does not leak any secret information).

REFERENCES

- Alejandro Aguirre, Gilles Barthe, Marco Gaboardi, Deepak Garg, and Pierre-Yves Strub. 2017. A relational logic for higher-order programs. *Proceedings of the ACM on Programming Languages* 1, ICFP (2017), 1–29.
- Torben Amtoft, Sruthi Bandhakavi, and Anindya Banerjee. 2006. A Logic for Information Flow in Object-Oriented Programs. *SIGPLAN Not.* 41, 1 (jan 2006), 91–102. <https://doi.org/10.1145/1111320.1111046>
- Timos Antonopoulos, Eric Koskinen, Ton Chanh Le, Ramana Nagasamudram, David A. Naumann, and Minh Ngo. 2023. An Algebra of Alignment for Relational Verification. *Proc. ACM Program. Lang.* 7, POPL, Article 20 (jan 2023), 31 pages. <https://doi.org/10.1145/3571213>
- Mounir Assaf, David A Naumann, Julien Signoles, Eric Totel, and Frédéric Tronel. 2017. Hypercollecting semantics and its application to static analysis of information flow. *ACM SIGPLAN Notices* 52, 1 (2017), 874–887.
- Gilles Barthe, Pedro R D’argenio, and Tamara Rezk. 2011. Secure information flow by self-composition. *Mathematical Structures in Computer Science* 21, 6 (2011), 1207–1252.
- Gilles Barthe, Renate Eilers, Pamina Georgiou, Bernhard Gleiss, Laura Kovács, and Matteo Maffei. 2019. Verifying relational properties using trace logic. In *2019 Formal Methods in Computer Aided Design (FMCAD)*. IEEE, 170–178.
- Nick Benton. 2004. Simple Relational Correctness Proofs for Static Analyses and Program Transformations. In *Proceedings of the 31st ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages* (Venice, Italy) (POPL ’04). Association for Computing Machinery, New York, NY, USA, 14–25. <https://doi.org/10.1145/964001.964003>
- Raven Beutner and Bernd Finkbeiner. 2022. Software Verification of Hyperproperties Beyond k-Safety. In *Computer Aided Verification*, Sharon Shoham and Yakir Vizel (Eds.). Cham, 341–362.
- Raven Beutner and Bernd Finkbeiner. 2023. AutoHyper: Explicit-State Model Checking for HyperLTL. In *Tools and Algorithms for the Construction and Analysis of Systems*, Sriram Sankaranarayanan and Natasha Sharygina (Eds.). Springer Nature Switzerland, Cham, 145–163.
- Nikolaj Bjørner, Arie Gurfinkel, Ken McMillan, and Andrey Rybalchenko. 2015. *Horn Clause Solvers for Program Verification*. Springer International Publishing, Cham, 24–51. https://doi.org/10.1007/978-3-319-23534-9_2
- Sam Blackshear, Nikos Gorogiannis, Peter W. O’Hearn, and Ilya Sergey. 2018. RacerD: Compositional Static Race Detection. *Proc. ACM Program. Lang.* 2, OOPSLA, Article 144 (oct 2018), 28 pages. <https://doi.org/10.1145/3276514>
- Roberto Bruni, Roberto Giacobazzi, Roberta Gori, and Francesco Ranzato. 2021. A Logic for Locally Complete Abstract Interpretations. In *2021 36th Annual ACM/IEEE Symposium on Logic in Computer Science (LICS)*. 1–13. <https://doi.org/10.1109/LICS52264.2021.9470608>
- Edmund M Clarke. 1997. Model checking. In *Foundations of Software Technology and Theoretical Computer Science: 17th Conference Kharagpur, India, December 18–20, 1997 Proceedings* 17. Springer, 54–56.
- Michael R Clarkson, Bernd Finkbeiner, Masoud Koleini, Kristopher K Micinski, Markus N Rabe, and César Sánchez. 2014. Temporal logics for hyperproperties. In *International Conference on Principles of Security and Trust*. 265–284.
- Michael R. Clarkson and Fred B. Schneider. 2008. Hyperproperties. In *21st IEEE Computer Security Foundations Symposium*. 51–65. <https://doi.org/10.1109/CSF.2008.7>
- Norine Coenen, Bernd Finkbeiner, César Sánchez, and Leander Tentrup. 2019. Verifying hyperliveness. In *International Conference on Computer Aided Verification*. 121–139.
- David Costanzo and Zhong Shao. 2014. A Separation Logic for Enforcing Declarative Information Flow Control Policies. In *Principles of Security and Trust*, Martin Abadi and Steve Kremer (Eds.). 179–198.
- Patrick Cousot and Radhia Cousot. 1977. Abstract interpretation: a unified lattice model for static analysis of programs by construction or approximation of fixpoints. In *Proceedings of the 4th ACM SIGACT-SIGPLAN symposium on Principles of programming languages*. 238–252.
- Thibault Dardinier and Peter Müller. 2023. Hyper Hoare Logic: (Dis-)Proving Program Hyperproperties (extended version). [arXiv:cs.LO/2301.10037](https://arxiv.org/abs/2301.10037)
- Leonardo de Moura and Nikolaj Bjørner. 2008. Z3: An Efficient SMT Solver. In *Tools and Algorithms for the Construction and Analysis of Systems*, C. R. Ramakrishnan and Jakob Rehof (Eds.). 337–340.
- Edsko de Vries and Vasileios Koutavas. 2011. Reverse Hoare Logic. In *Software Engineering and Formal Methods*, Gilles Barthe, Alberto Pardo, and Gerardo Schneider (Eds.). 155–171.
- Robert Dickerson, Qianchuan Ye, Michael K. Zhang, and Benjamin Delaware. 2022. RHLE: Modular Deductive Verification of Relational $\forall\exists$ Properties. In *Programming Languages and Systems: 20th Asian Symposium, APLAS 2022, Auckland, New Zealand, December 5, 2022, Proceedings* (Auckland, New Zealand). 67–87. https://doi.org/10.1007/978-3-031-21037-2_4
- Dino Distefano, Manuel Fähndrich, Francesco Logozzo, and Peter W. O’Hearn. 2019. Scaling Static Analyses at Facebook. *Commun. ACM* 62, 8 (jul 2019), 62–70. <https://doi.org/10.1145/3338112>
- Emanuele D’Osualdo, Azadeh Farzan, and Derek Dreyer. 2022. Proving Hypersafety Compositionally. *Proc. ACM Program. Lang.* 6, OOPSLA2, Article 135 (oct 2022), 26 pages. <https://doi.org/10.1145/3563298>
- Marco Eilers, Thibault Dardinier, and Peter Müller. 2023. CommCSL: Proving Information Flow Security for Concurrent Programs Using Abstract Commutativity. *Proc. ACM Program. Lang.* 7, PLDI, Article 175 (jun 2023), 26 pages. <https://doi.org/10.1145/3563298>

- [//doi.org/10.1145/3591289](https://doi.org/10.1145/3591289)
- Marco Eilers, Peter Müller, and Samuel Hitz. 2019. Modular product programs. *ACM Transactions on Programming Languages and Systems (TOPLAS)* 42, 1 (2019), 1–37.
- Gidon Ernst and Toby Murray. 2019. SecCSL: Security Concurrent Separation Logic. In *Computer Aided Verification*, Isil Dillig and Serdar Tasiran (Eds.). Cham, 208–230.
- Azadeh Farzan and Anthony Vandikas. 2019. Automated Hypersafety Verification. In *Computer Aided Verification*, Isil Dillig and Serdar Tasiran (Eds.). Springer International Publishing, Cham, 200–218.
- Jean-Christophe Filliâtre and Andrei Paskevich. 2013. Why3—where programs meet provers. In *Programming Languages and Systems: 22nd European Symposium on Programming, ESOP 2013, Held as Part of the European Joint Conferences on Theory and Practice of Software, ETAPS 2013, Rome, Italy, March 16–24, 2013. Proceedings 22*. Springer, 125–128.
- Bernd Finkbeiner, Markus N Rabe, and César Sánchez. 2015. Algorithms for model checking HyperLTL and HyperCTL. In *International Conference on Computer Aided Verification*. Springer, 30–48.
- Robert W. Floyd. 1967. Assigning Meanings to Programs. *Proceedings of Symposium in Applied Mathematics* (1967), 19–32.
- Nikos Gorogiannis, Peter W. O’Hearn, and Ilya Sergey. 2019. A True Positives Theorem for a Static Race Detector. *Proc. ACM Program. Lang.* 3, POPL, Article 57 (jan 2019), 29 pages. <https://doi.org/10.1145/3290370>
- David Harel. 1979. *First-order dynamic logic*. Springer.
- C. A. R. Hoare. 1969. An Axiomatic Basis for Computer Programming. *Commun. ACM* 12, 10 (oct 1969), 576–580. <https://doi.org/10.1145/363235.363259>
- Tzu-Han Hsu, César Sánchez, and Borzoo Bonakdarpour. 2021. Bounded Model Checking for Hyperproperties. In *Tools and Algorithms for the Construction and Analysis of Systems*, Jan Friso Groote and Kim Guldstrand Larsen (Eds.). Springer International Publishing, Cham, 94–112.
- Shachar Itzhaky, Sharon Shoham, and Yakir Vizel. 2024. Hyperproperty Verification as CHC Satisfiability. arXiv:cs.LO/2304.12588
- Ioannis T. Kassios. 2006. Dynamic Frames: Support for Framing, Dependencies and Sharing Without Restrictions. In *FM 2006: Formal Methods*, Jayadev Misra, Tobias Nipkow, and Emil Sekerinski (Eds.). Springer Berlin Heidelberg, Berlin, Heidelberg, 268–283.
- Dexter Kozen. 1997. Kleene algebra with tests. *ACM Trans. Program. Lang. Syst.* 19, 3 (may 1997), 427–443. <https://doi.org/10.1145/256167.256195>
- Quang Loc Le, Azalea Raad, Jules Villard, Josh Berdine, Derek Dreyer, and Peter W. O’Hearn. 2022. Finding Real Bugs in Big Programs with Incorrectness Logic. *Proc. ACM Program. Lang.* 6, OOPSLA1, Article 81 (apr 2022), 27 pages. <https://doi.org/10.1145/3527325>
- K. Rustan M. Leino. 2008. This is Boogie 2. (June 2008). <https://www.microsoft.com/en-us/research/publication/this-is-boogie-2-2/>
- K Rustan M Leino. 2010. Dafny: An automatic program verifier for functional correctness. In *International conference on logic for programming artificial intelligence and reasoning*. Springer, 348–370.
- K. Rustan M. Leino and Rosemary Monahan. 2009. Reasoning about comprehensions with first-order SMT solvers. In *Proceedings of the 2009 ACM Symposium on Applied Computing (Honolulu, Hawaii) (SAC ’09)*. Association for Computing Machinery, New York, NY, USA, 615–622. <https://doi.org/10.1145/1529282.1529411>
- Kenji Maillard, Cătălin Hrițcu, Exequiel Rivas, and Antoine Van Muylder. 2019. The next 700 Relational Program Logics. *Proc. ACM Program. Lang.* 4, POPL, Article 4 (dec 2019), 33 pages. <https://doi.org/10.1145/3371072>
- Petar Maksimović, Caroline Cronjäger, Andreas Löw, Julian Sutherland, and Philippa Gardner. 2023. Exact Separation Logic: Towards Bridging the Gap Between Verification and Bug-Finding. In *37th European Conference on Object-Oriented Programming (ECOOP 2023)*, Vol. 263. 19:1–19:27. <https://doi.org/10.4230/LIPIcs.ECOOP.2023.19>
- Daryl McCullough. 1987. Specifications for multi-level security and a hook-up. In *1987 IEEE Symposium on Security and Privacy*. IEEE, 161–161.
- John McLean. 1996. A general theory of composition for a class of “possibilistic” properties. *IEEE Transactions on Software Engineering* 22, 1 (1996), 53–67.
- P. Müller, M. Schwerhoff, and A. J. Summers. 2016. Viper: A Verification Infrastructure for Permission-Based Reasoning. In *Verification, Model Checking, and Abstract Interpretation (VMCAI) (LNCS)*, B. Jobstmann and K. R. M. Leino (Eds.), Vol. 9583. Springer-Verlag, 41–62. https://doi.org/10.1007/978-3-662-49122-5_2
- Toby Murray. 2020. An Under-Approximate Relational Logic: Heralding Logics of Insecurity, Incorrect Implementation and More. <https://doi.org/10.48550/ARXIV.2003.04791>
- Ramana Nagasamudram, Anindya Banerjee, and David A. Naumann. 2023. The WhyRel Prototype for Modular Relational Verification of Pointer Programs. In *Tools and Algorithms for the Construction and Analysis of Systems*, Sriram Sankaranarayanan and Natasha Sharygina (Eds.). Springer Nature Switzerland, Cham, 133–151.
- David A. Naumann and Minh Ngo. 2019. Whither Specifications as Programs. In *Unifying Theories of Programming*, Pedro Ribeiro and Augusto Sampaio (Eds.). Springer International Publishing, Cham, 39–61.

- Tobias Nipkow, Markus Wenzel, and Lawrence C. Paulson. 2002. *Isabelle/HOL: A Proof Assistant for Higher-Order Logic*. Springer-Verlag, Berlin, Heidelberg.
- Peter W. O'Hearn. 2019. Incorrectness Logic. *Proc. ACM Program. Lang.* 4, POPL, Article 10 (dec 2019), 32 pages. <https://doi.org/10.1145/3371078>
- J.C. Reynolds. 2002. Separation logic: a logic for shared mutable data structures. In *Proceedings 17th Annual IEEE Symposium on Logic in Computer Science*. 55–74. <https://doi.org/10.1109/LICS.2002.1029817>
- Marcelo Sousa and Isil Dillig. 2016. Cartesian Hoare Logic for Verifying K-Safety Properties. In *Proceedings of the 37th ACM SIGPLAN Conference on Programming Language Design and Implementation* (Santa Barbara, CA, USA) (PLDI '16). Association for Computing Machinery, New York, NY, USA, 57–69. <https://doi.org/10.1145/2908080.2908092>
- Tachio Terauchi and Alex Aiken. 2005. Secure information flow as a safety problem. In *International Static Analysis Symposium*. 352–367.
- Hiroshi Unno, Tachio Terauchi, and Eric Koskinen. 2021. Constraint-Based Relational Verification. In *Computer Aided Verification*, Alexandra Silva and K. Rustan M. Leino (Eds.). Springer International Publishing, Cham, 742–766.
- Dennis Volpano, Cynthia Irvine, and Geoffrey Smith. 1996. A sound type system for secure flow analysis. *Journal of computer security* 4, 2-3 (1996), 167–187.
- D. Volpano and G. Smith. 1997. Eliminating covert flows with minimum typings. In *Proceedings 10th Computer Security Foundations Workshop*. 156–168. <https://doi.org/10.1109/CSFW.1997.596807>
- Hongseok Yang. 2007. Relational separation logic. *Theoretical Computer Science* 375, 1 (2007), 308–334. <https://doi.org/10.1016/j.tcs.2006.12.036>
- Noam Zilberstein, Derek Dreyer, and Alexandra Silva. 2023. Outcome Logic: A Unifying Foundation for Correctness and Incorrectness Reasoning. <https://www.cs.cornell.edu/~noamz/files/pubs/outcome.pdf>